

Data Visualization Group

Northstar Middle School

Denny Lee

About Me

- Sr Staff Developer Advocate, Databricks
- Senior Director of Data Science Engineering at SAP Concur
- Principal Program Manager at Microsoft for Azure Cosmos DB, Project Isotope (Azure HDInsight), SQL Server



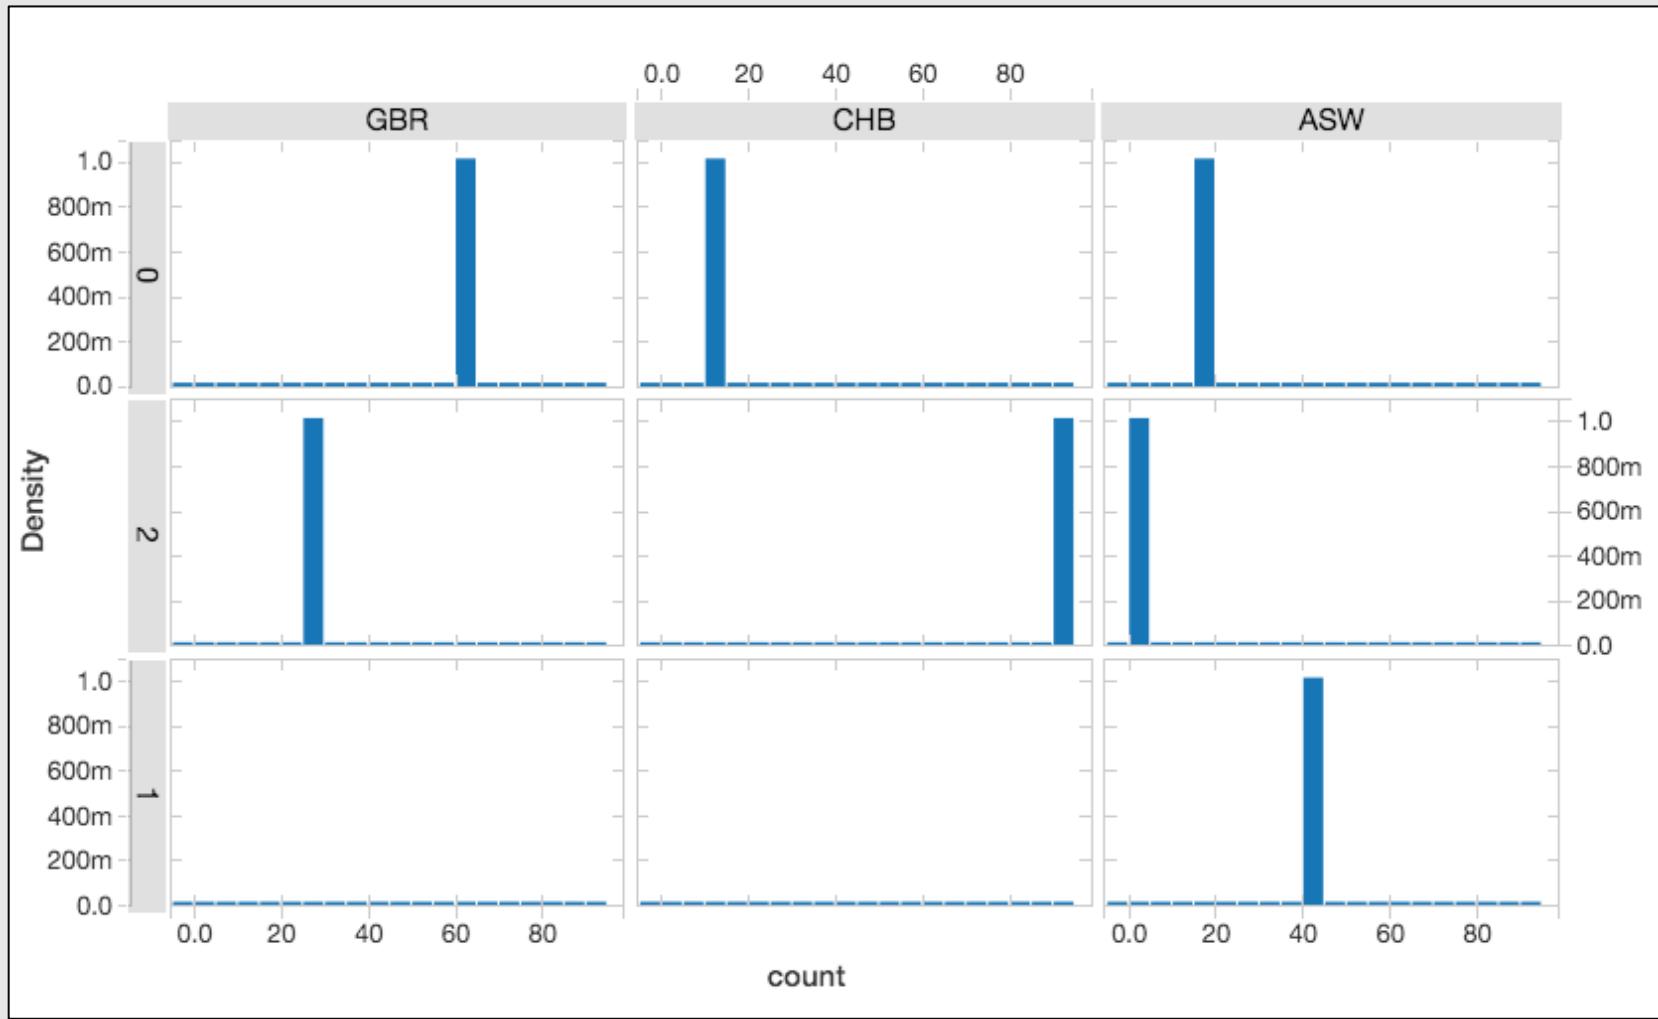
Denny surprised that he's still awake

Session Goals

- What is Deep Learning?
- Deep Learning Applications – including data visualization
- Biology vs scientific model
- Demos!
- Maths (completely optional)

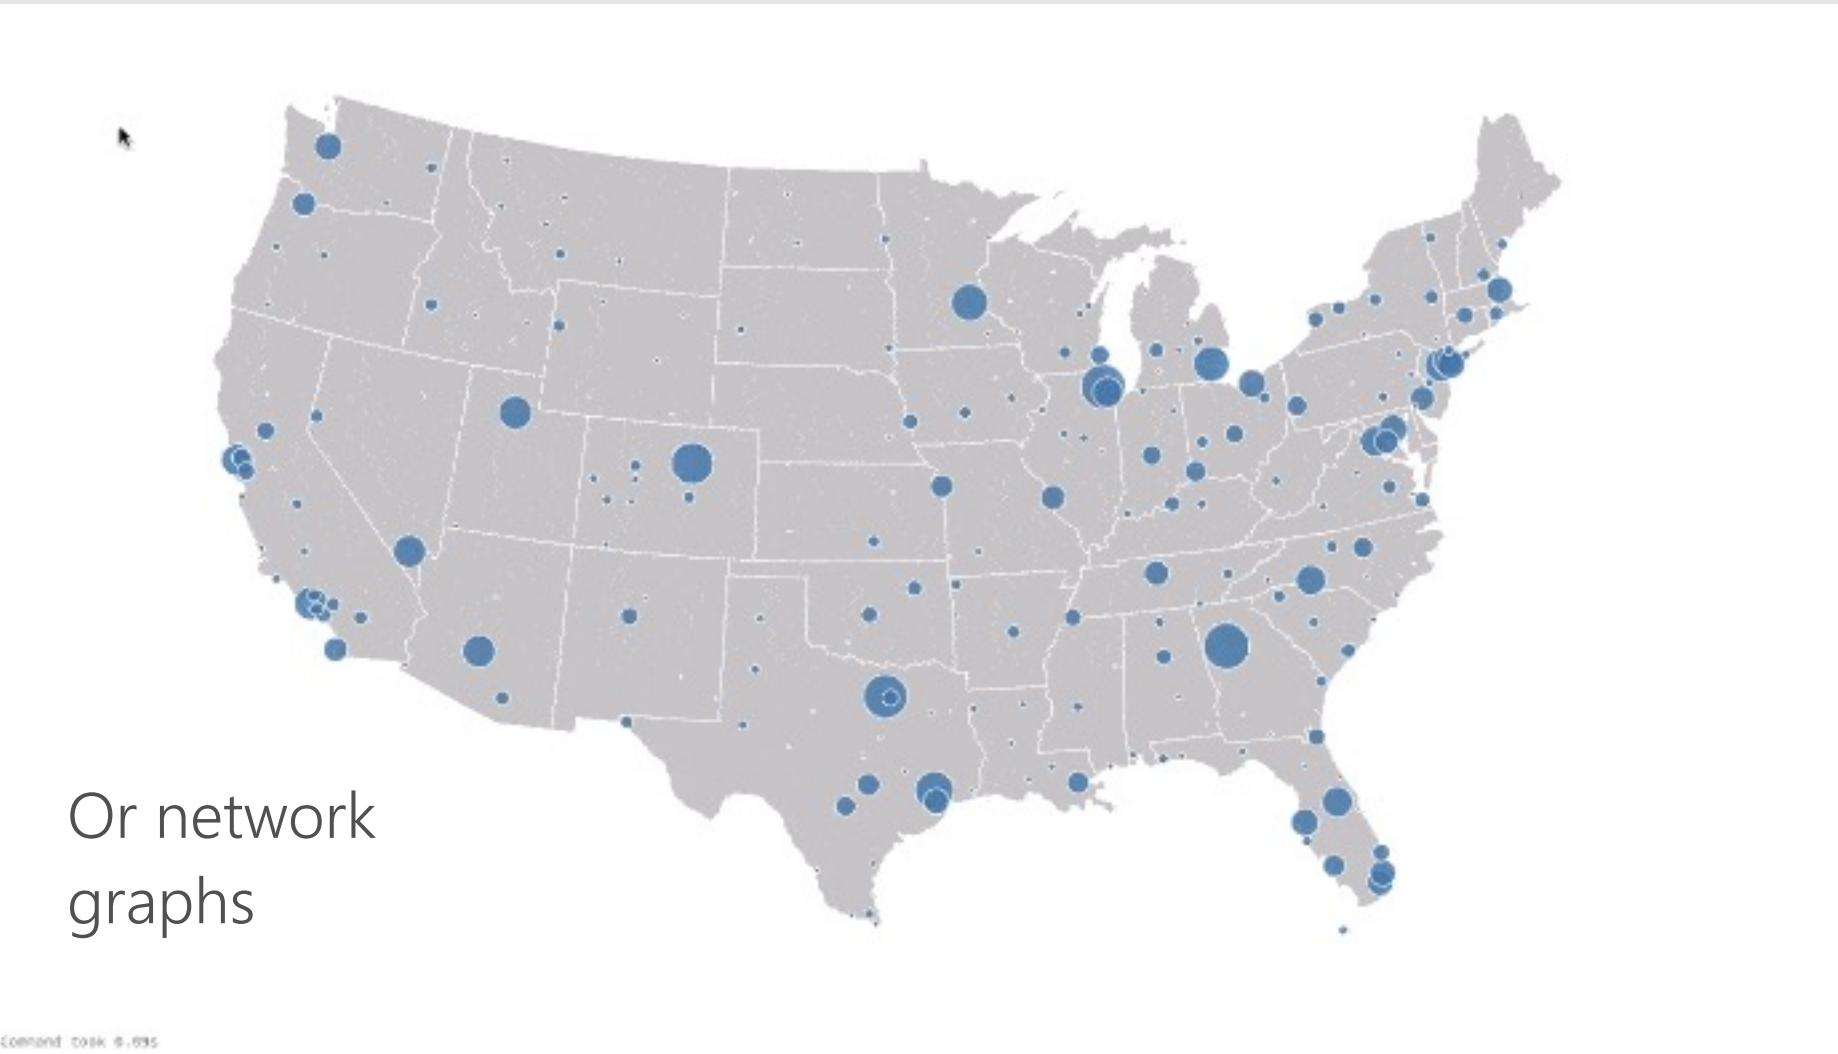
What is Deep Learning?

**WAIT, ISN'T THIS
A VISUALIZATION SESSION?**



Data visualization
can mean a lot of
things

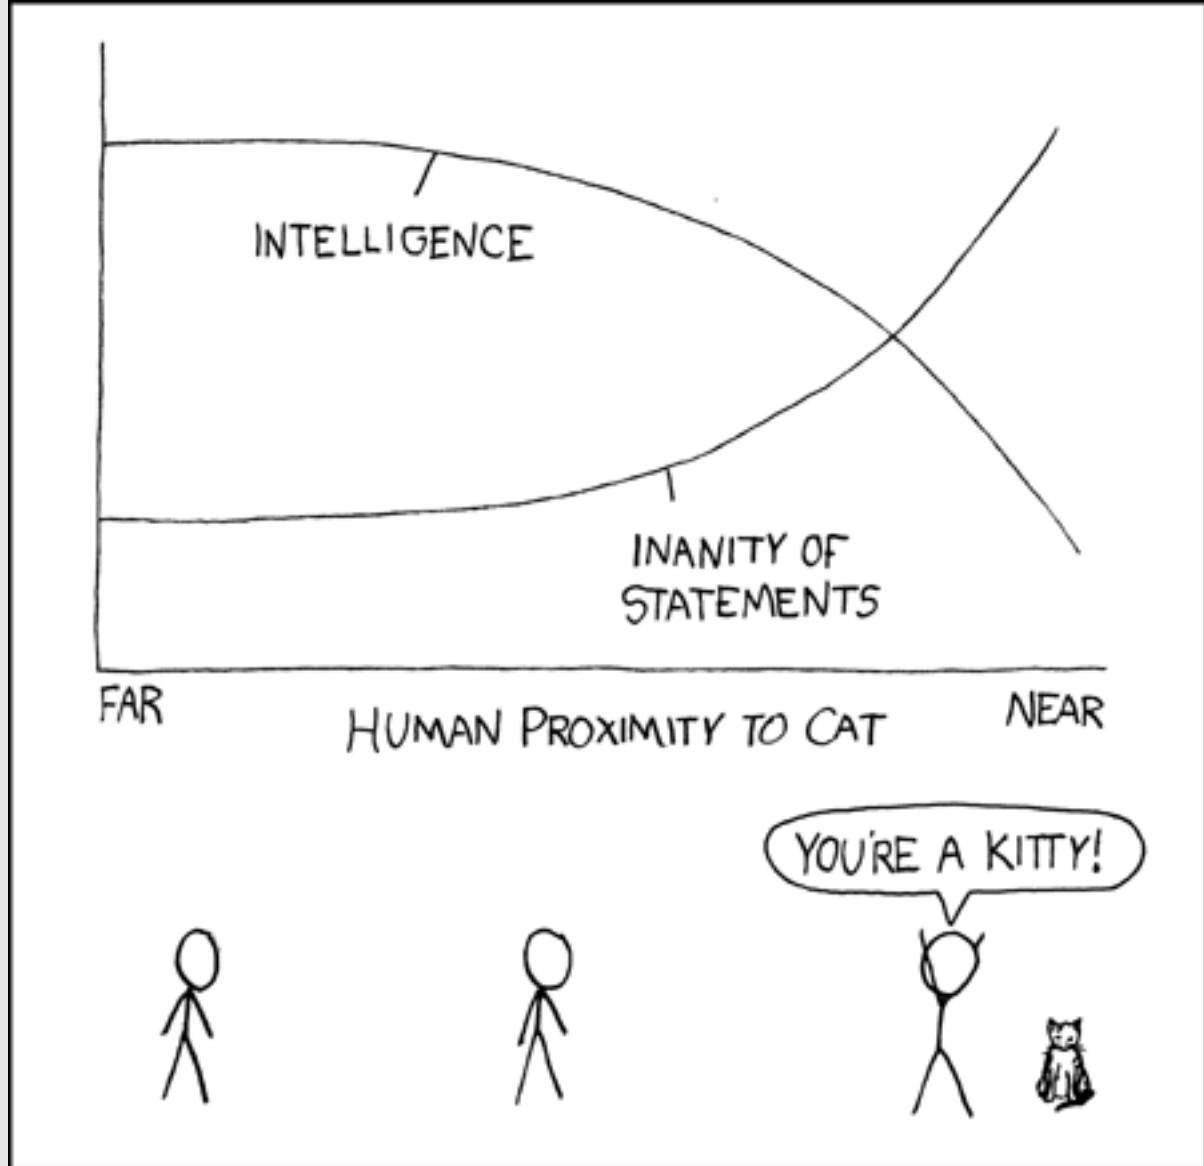
.. like graphs!



Or network
graphs

Command took 0.695

Or even more charts ...



[xkcd comic 231](#)

Or even worse ...



But today,
we're talking
about the
ability to
interpret
visual data

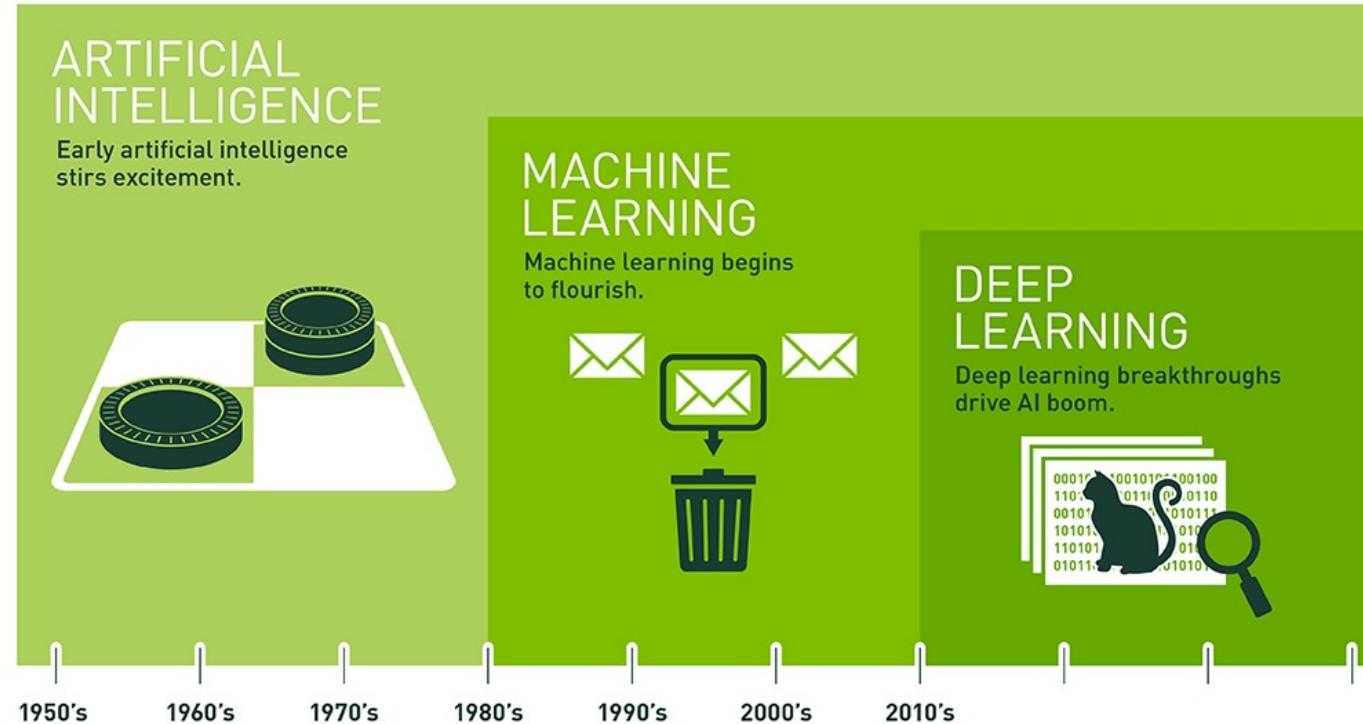


source: <http://bit.ly/2jcuuh5>

Deep Learning

The next step in AI

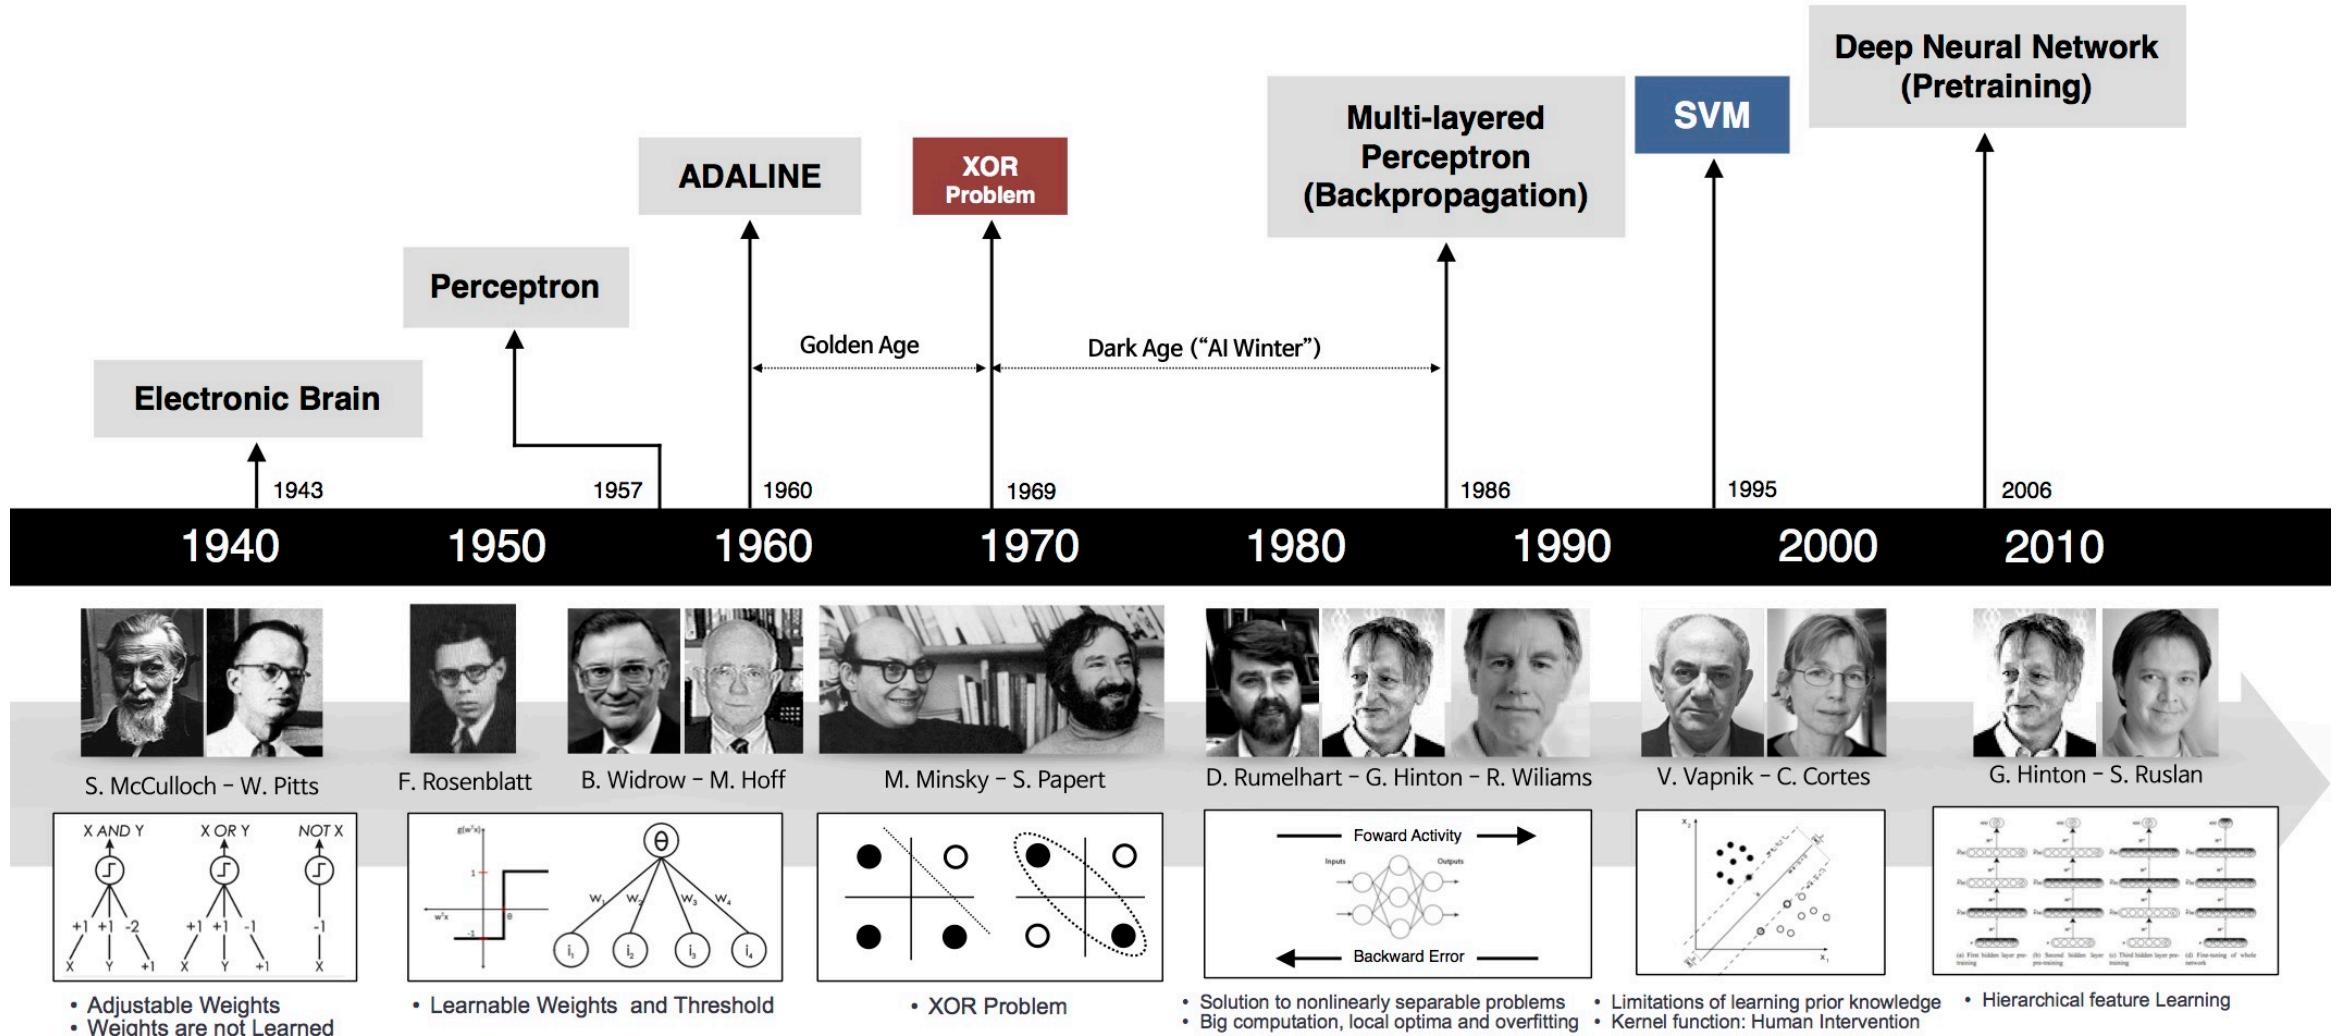
Learning... Lots of learning...



Since an early flush of optimism in the 1950s, smaller subsets of artificial intelligence – first machine learning, then deep learning, a subset of machine learning – have created ever larger disruptions.

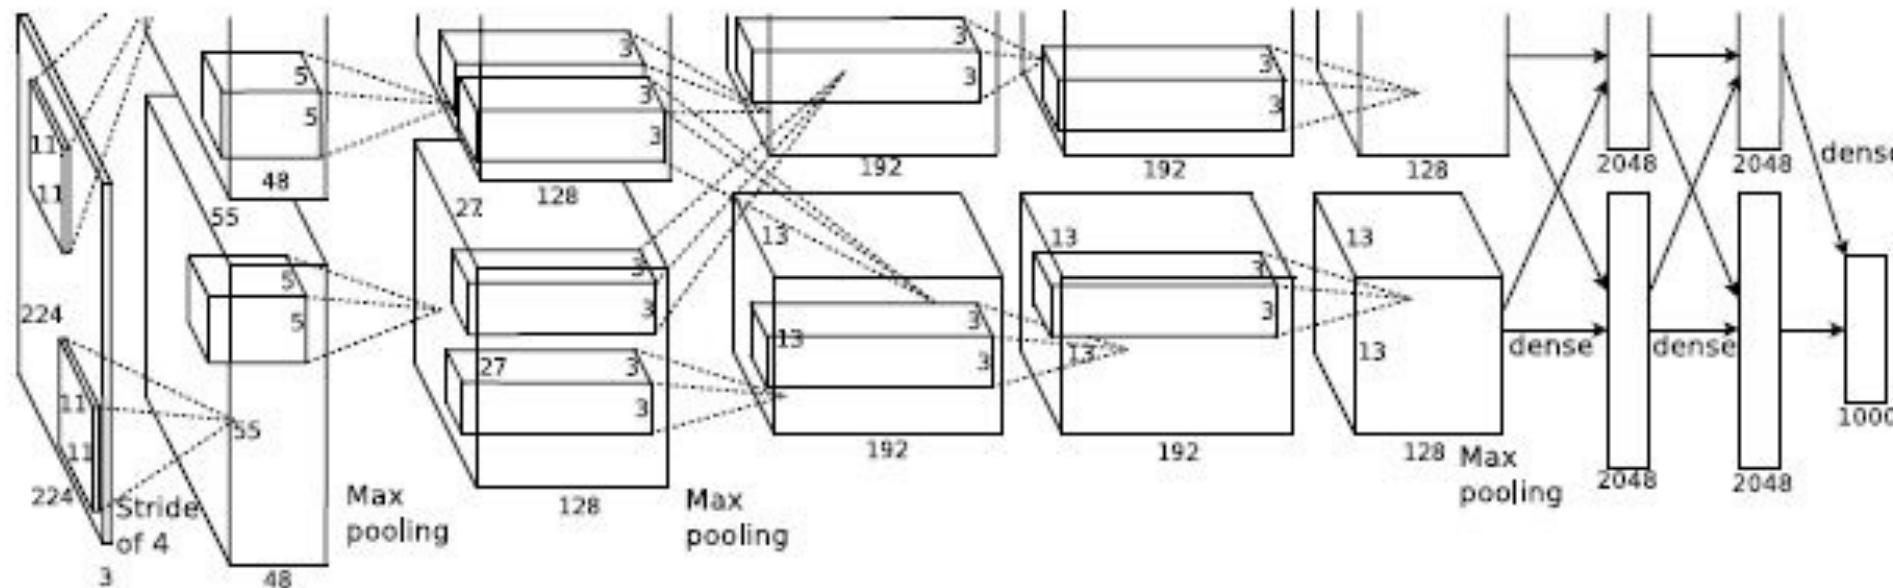
Source: <http://bit.ly/2suufGJ>

History of AI research



Deep learning

- Vast and sophisticated ANN structure
- Tens, hundreds, thousands... of hidden layers
- Ability to not only map the input to the output but recognize what inputs are needed for the task

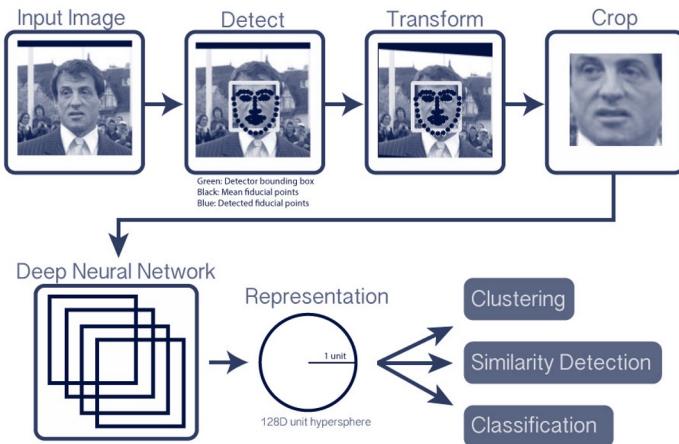


Example: AlexNet (ImageNet 2012 winner); Source: <http://bit.ly/2tNa3Ps>

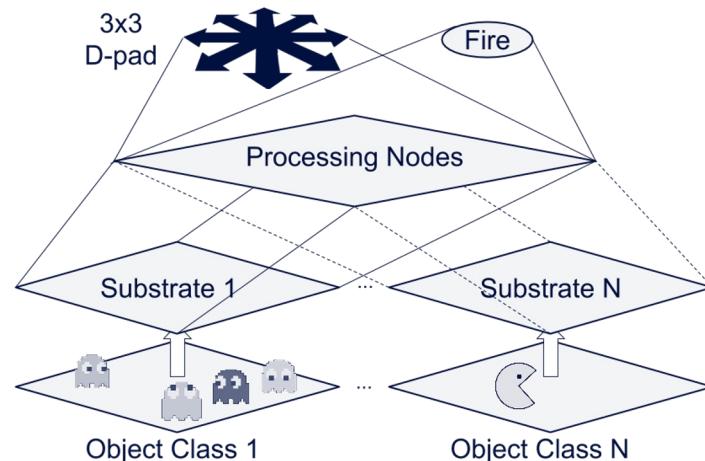
Deep Learning Applications

Applications of neural networks

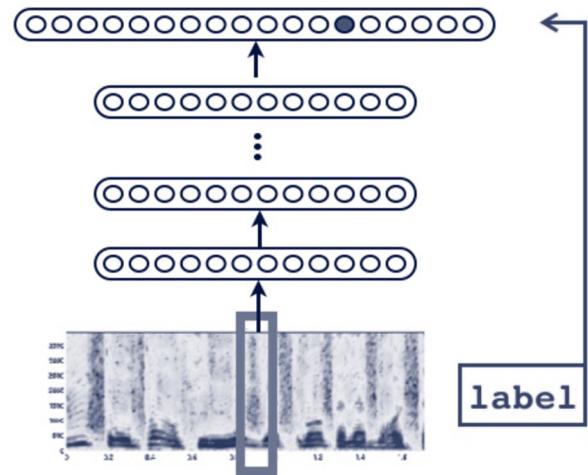
Object Recognition: Facial



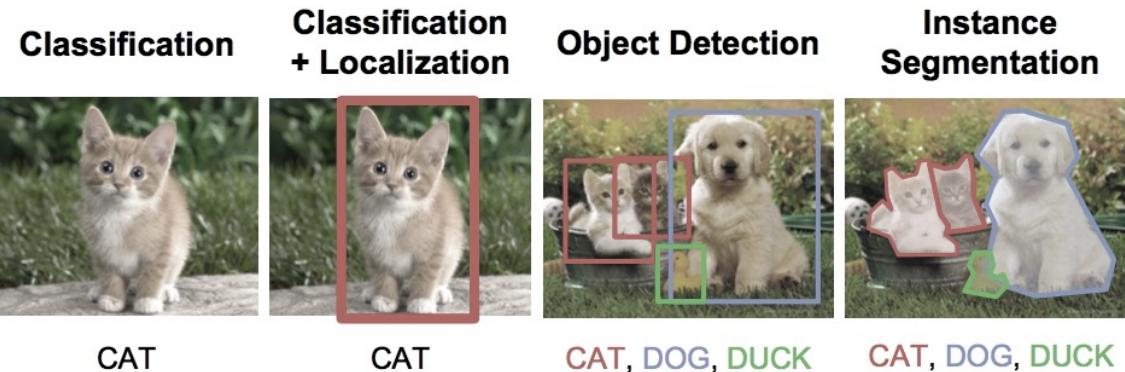
Game Playing



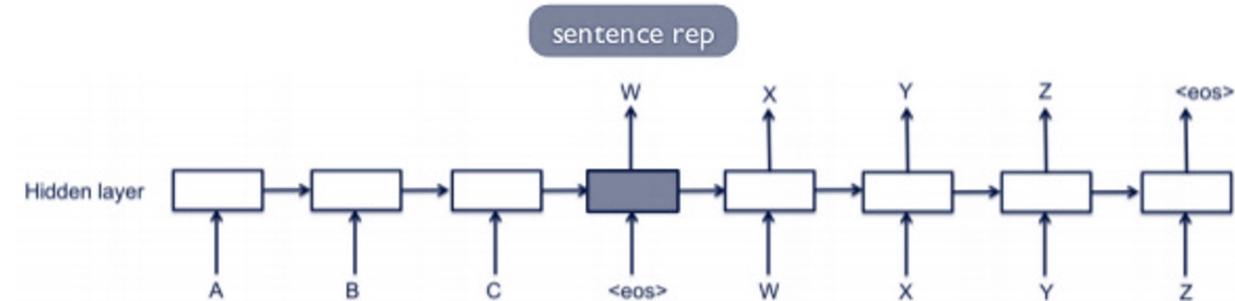
Object Recognition: Speech



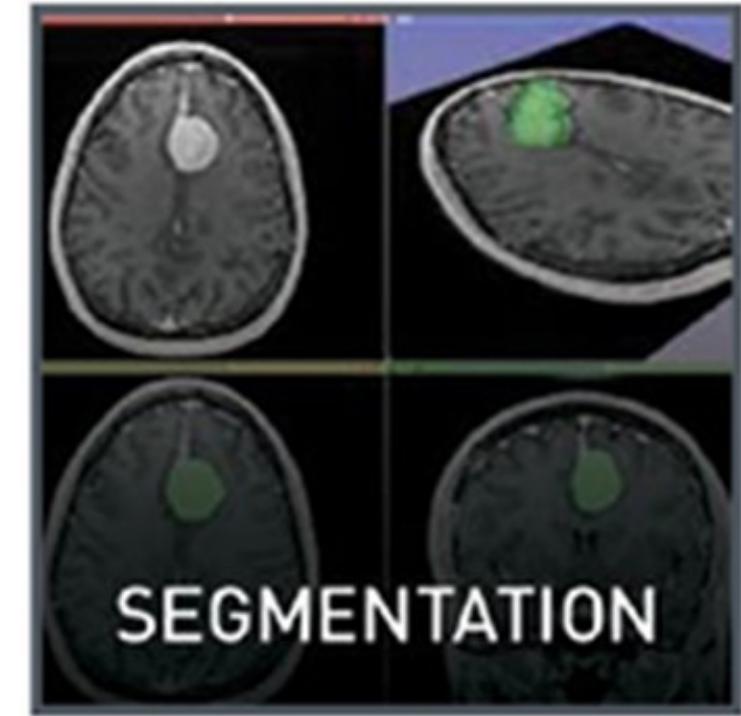
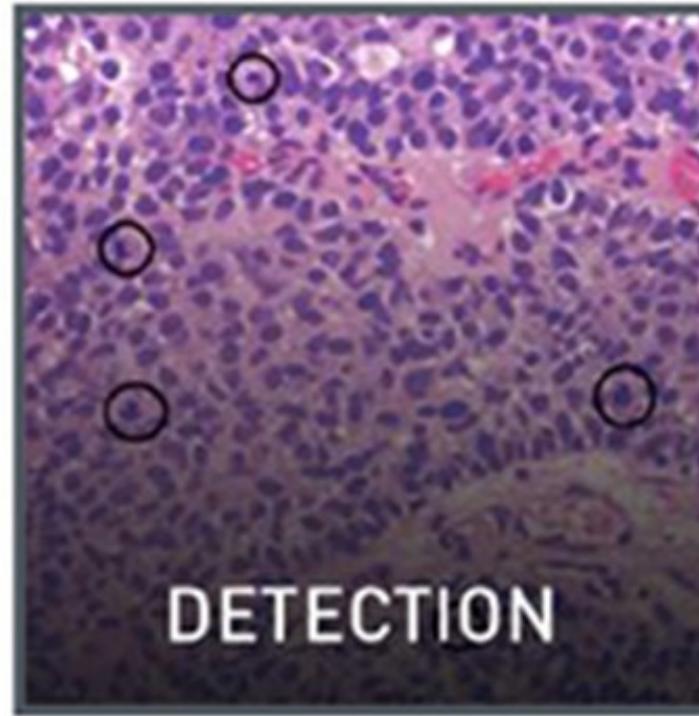
Object Classification



Language Translation

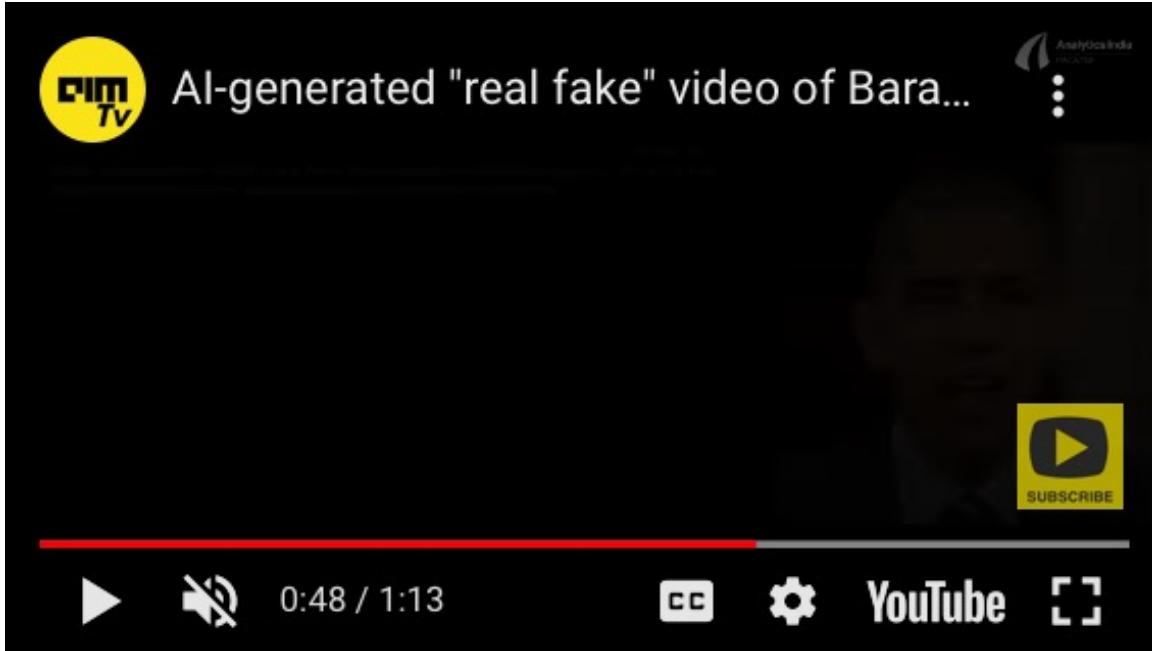


Healthcare



Source: <http://bit.ly/2BXF5sR>

Faking videos



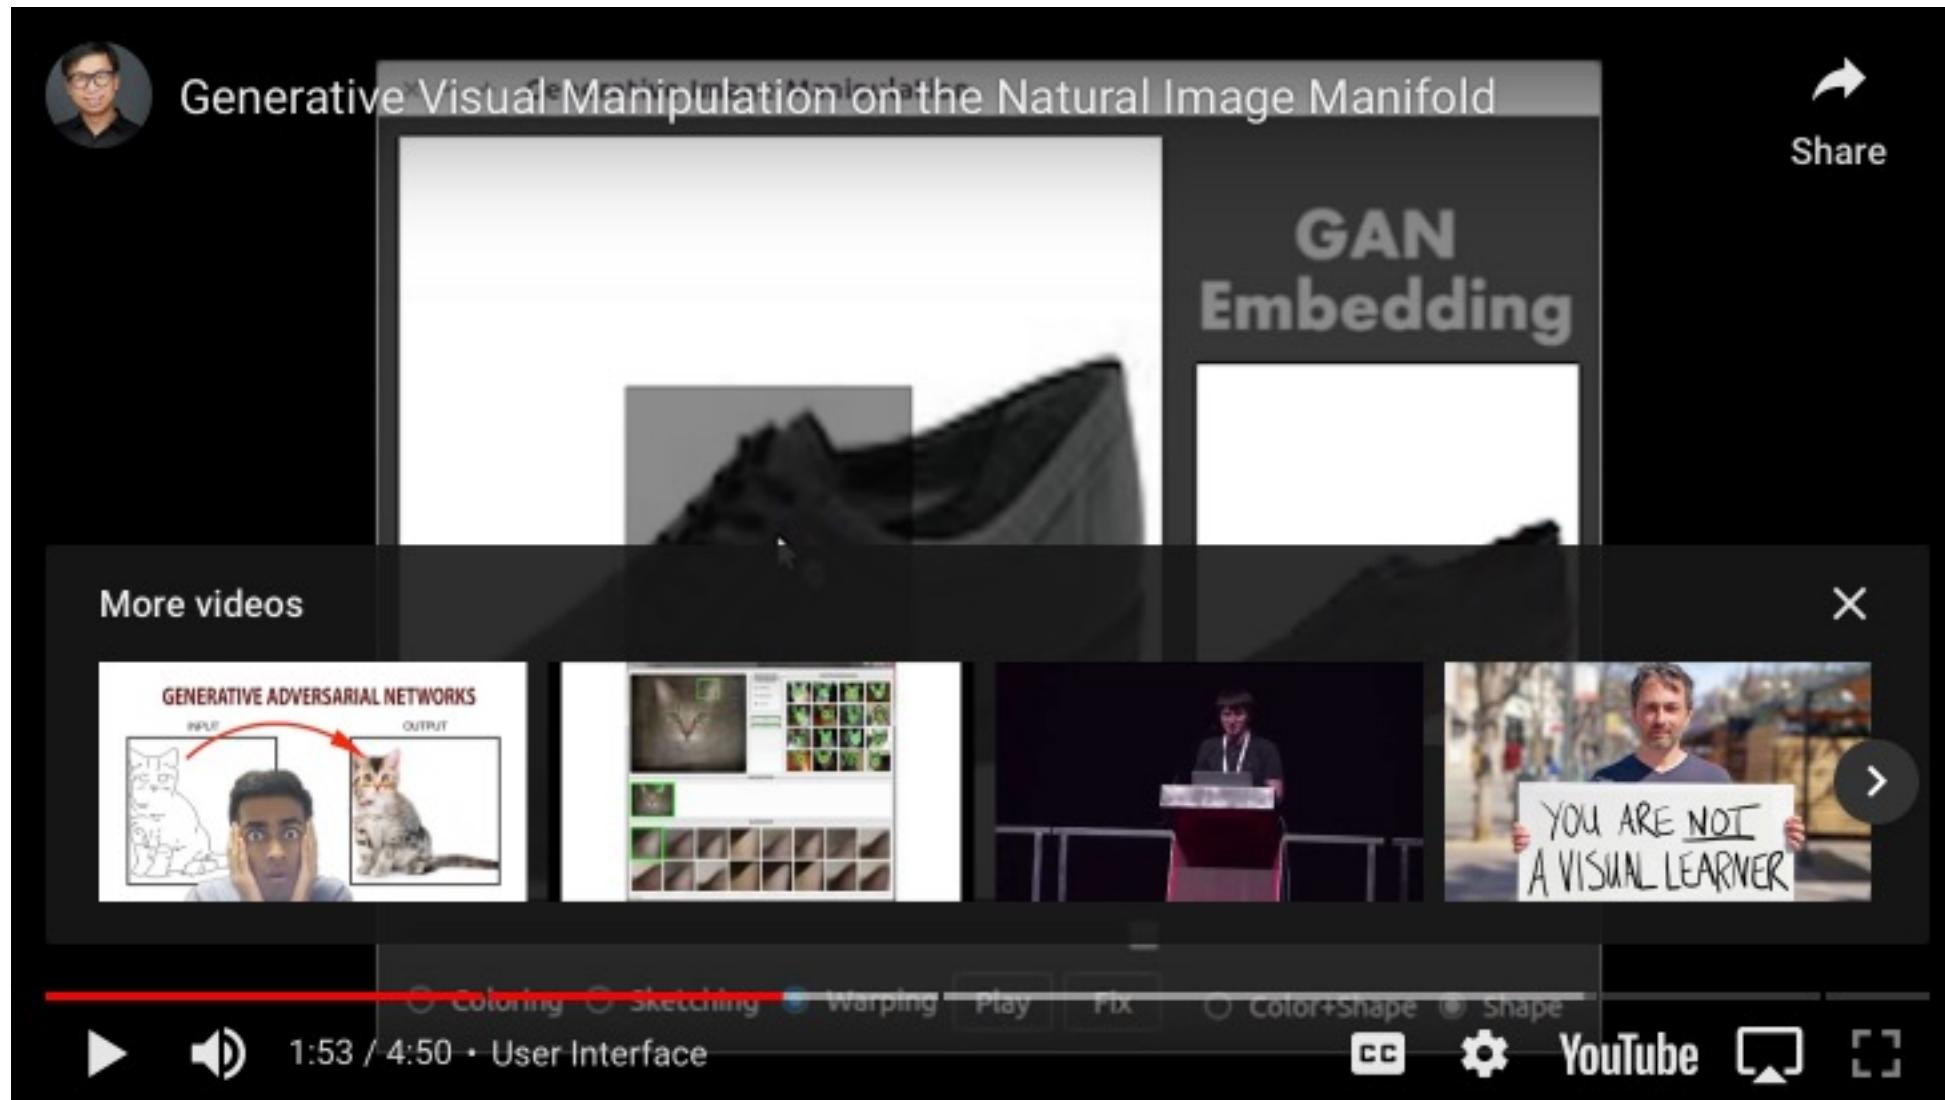
Source: <https://youtu.be/dkoi7sZvWiU>



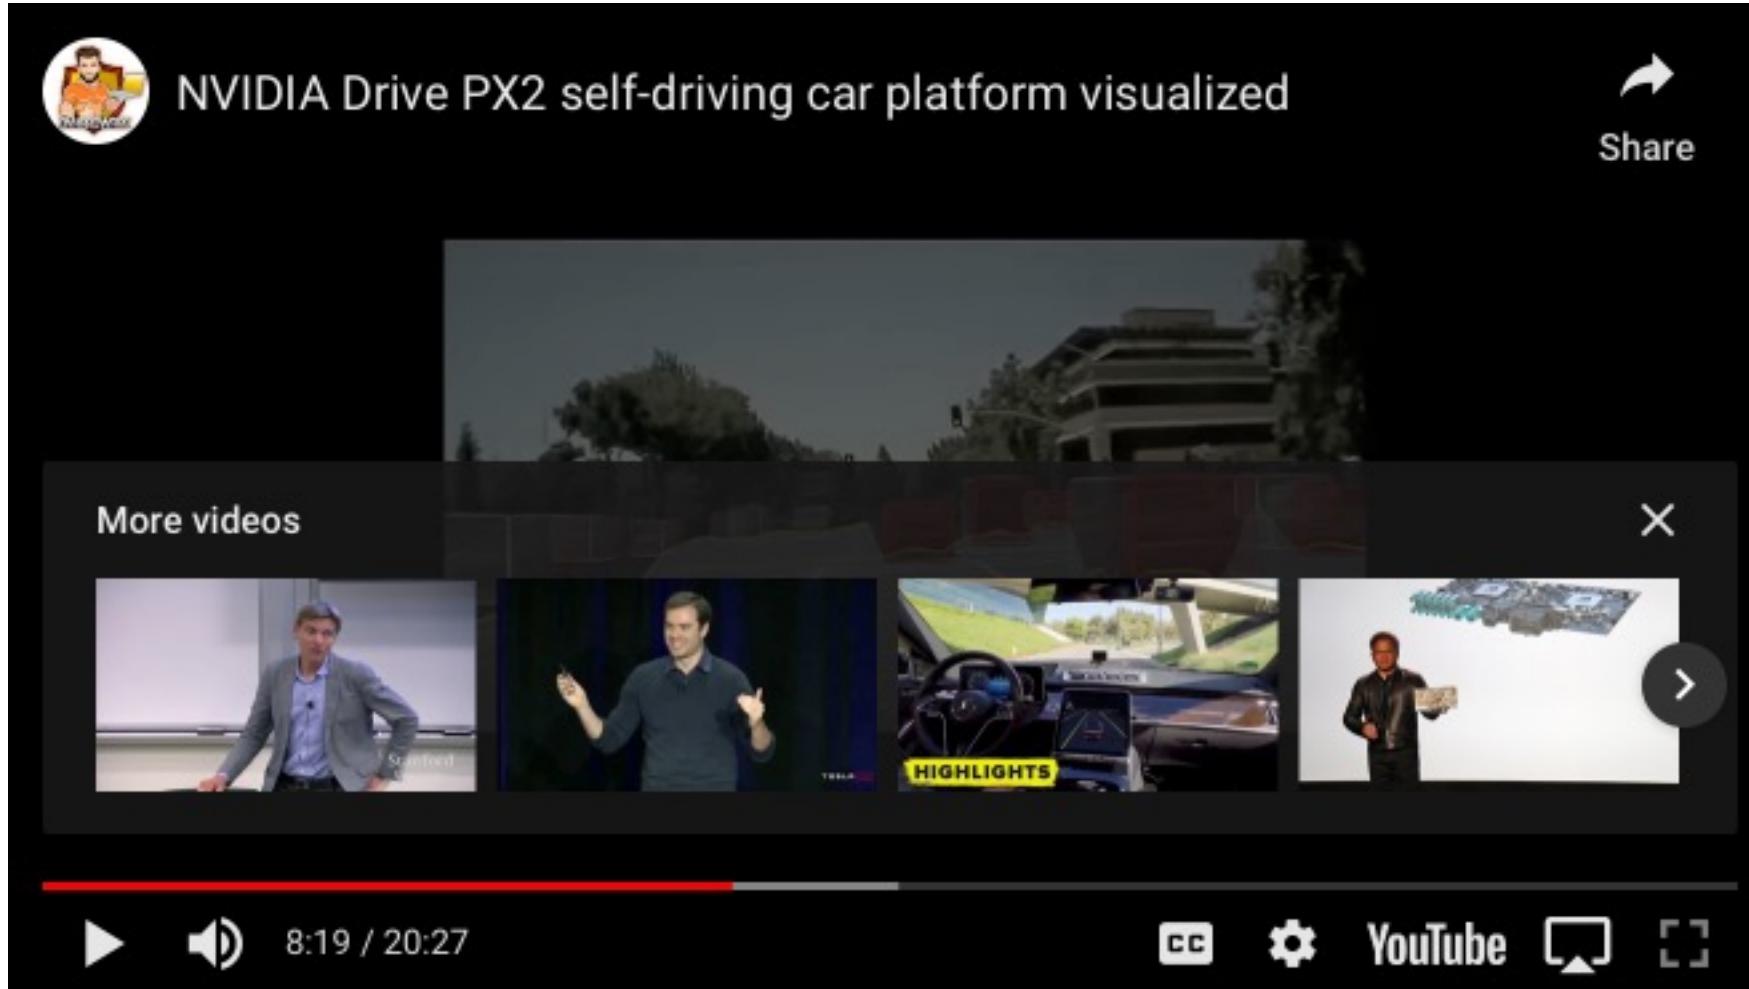
Source: <https://youtu.be/JzgOfISLNjk>

Drawing support

<https://www.youtube.com/watch?v=9c4z6YsBGQ0>

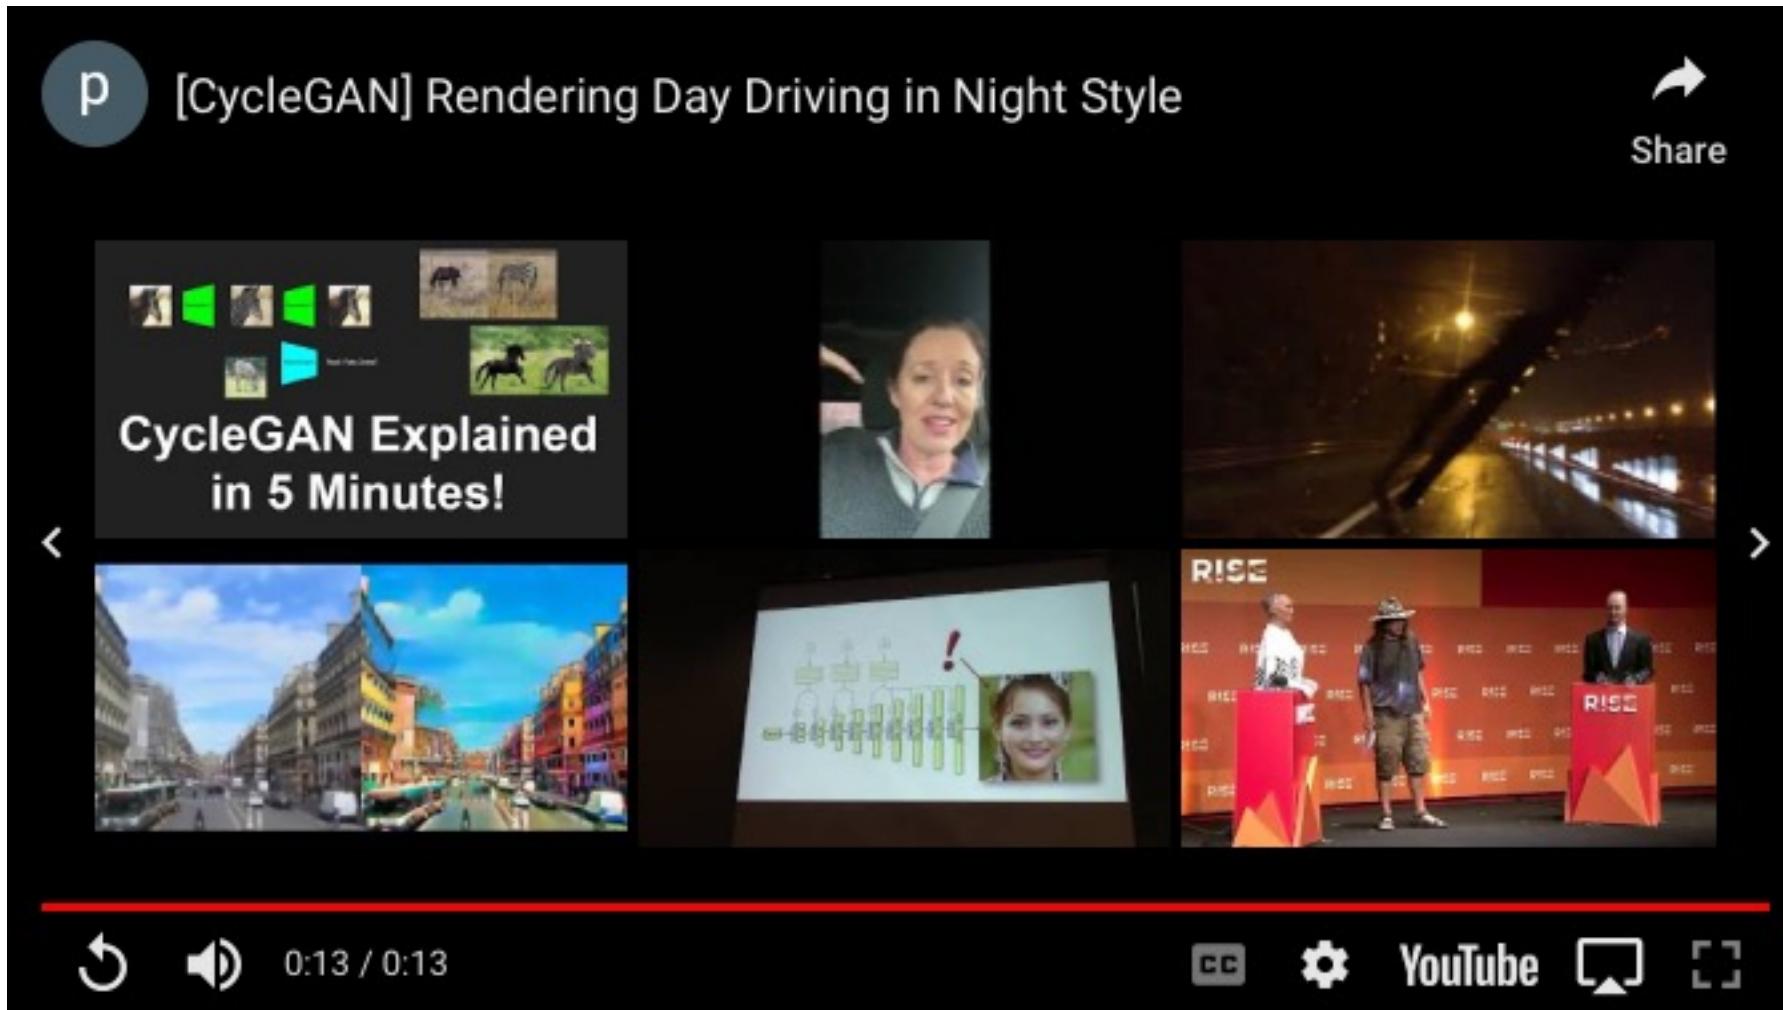


Self-driving cars



Source: <https://youtu.be/URmxzxYlmtg?t=6m30s>

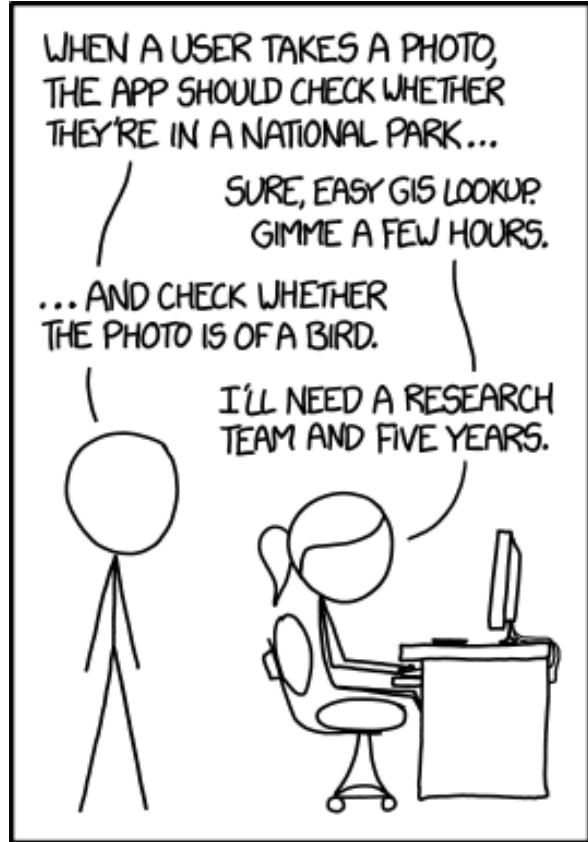
Turning the day into night



Source: <https://youtu.be/N7KbfWodXJE>

Recognizing Objects

xkcd #1425 ([View original here](#))



Really good explanation of convolution
<http://bit.ly/2suZGkk>

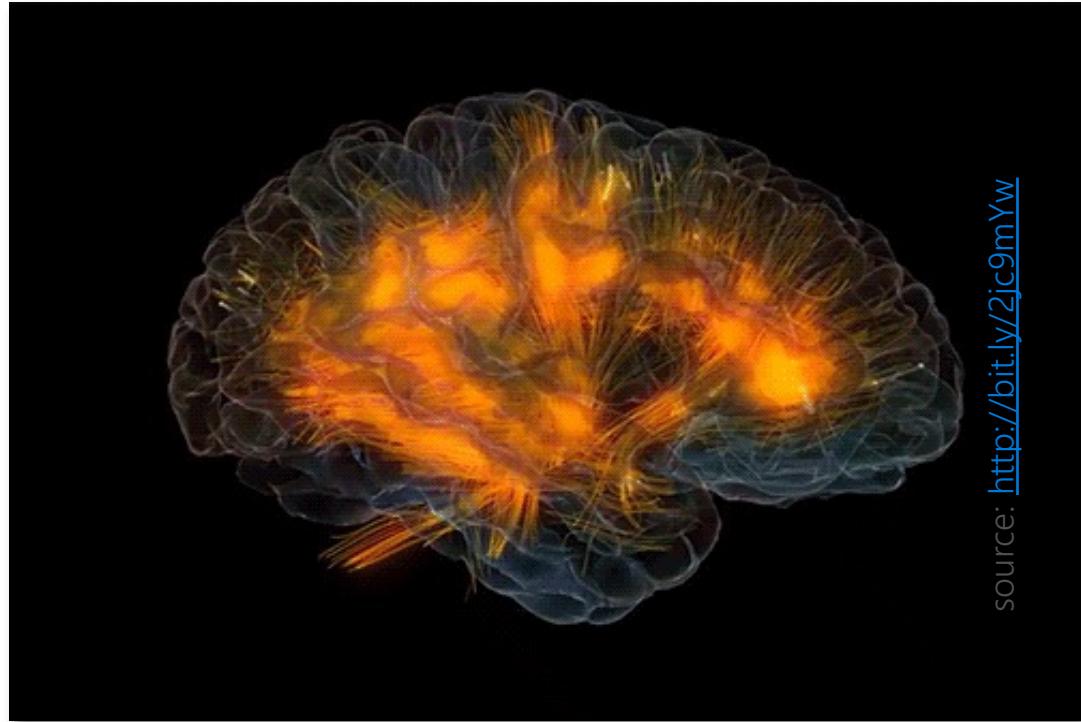


source: <http://bit.ly/2jjzFyt>

What are artificial neural networks?

An imitation of human's brain

- Dense net of simple structures
- Around 100 billion neurons
- Each connected to ~10k other neurons
- 10^{15} synaptic connections



source: <http://bit.ly/2ic9mYw>

A neuron at work

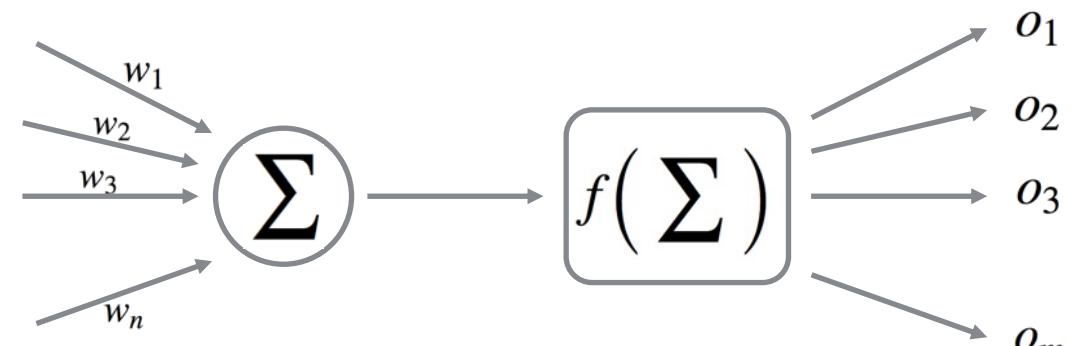
- Receive signal
- Process it
- Propagate the signal (or not)



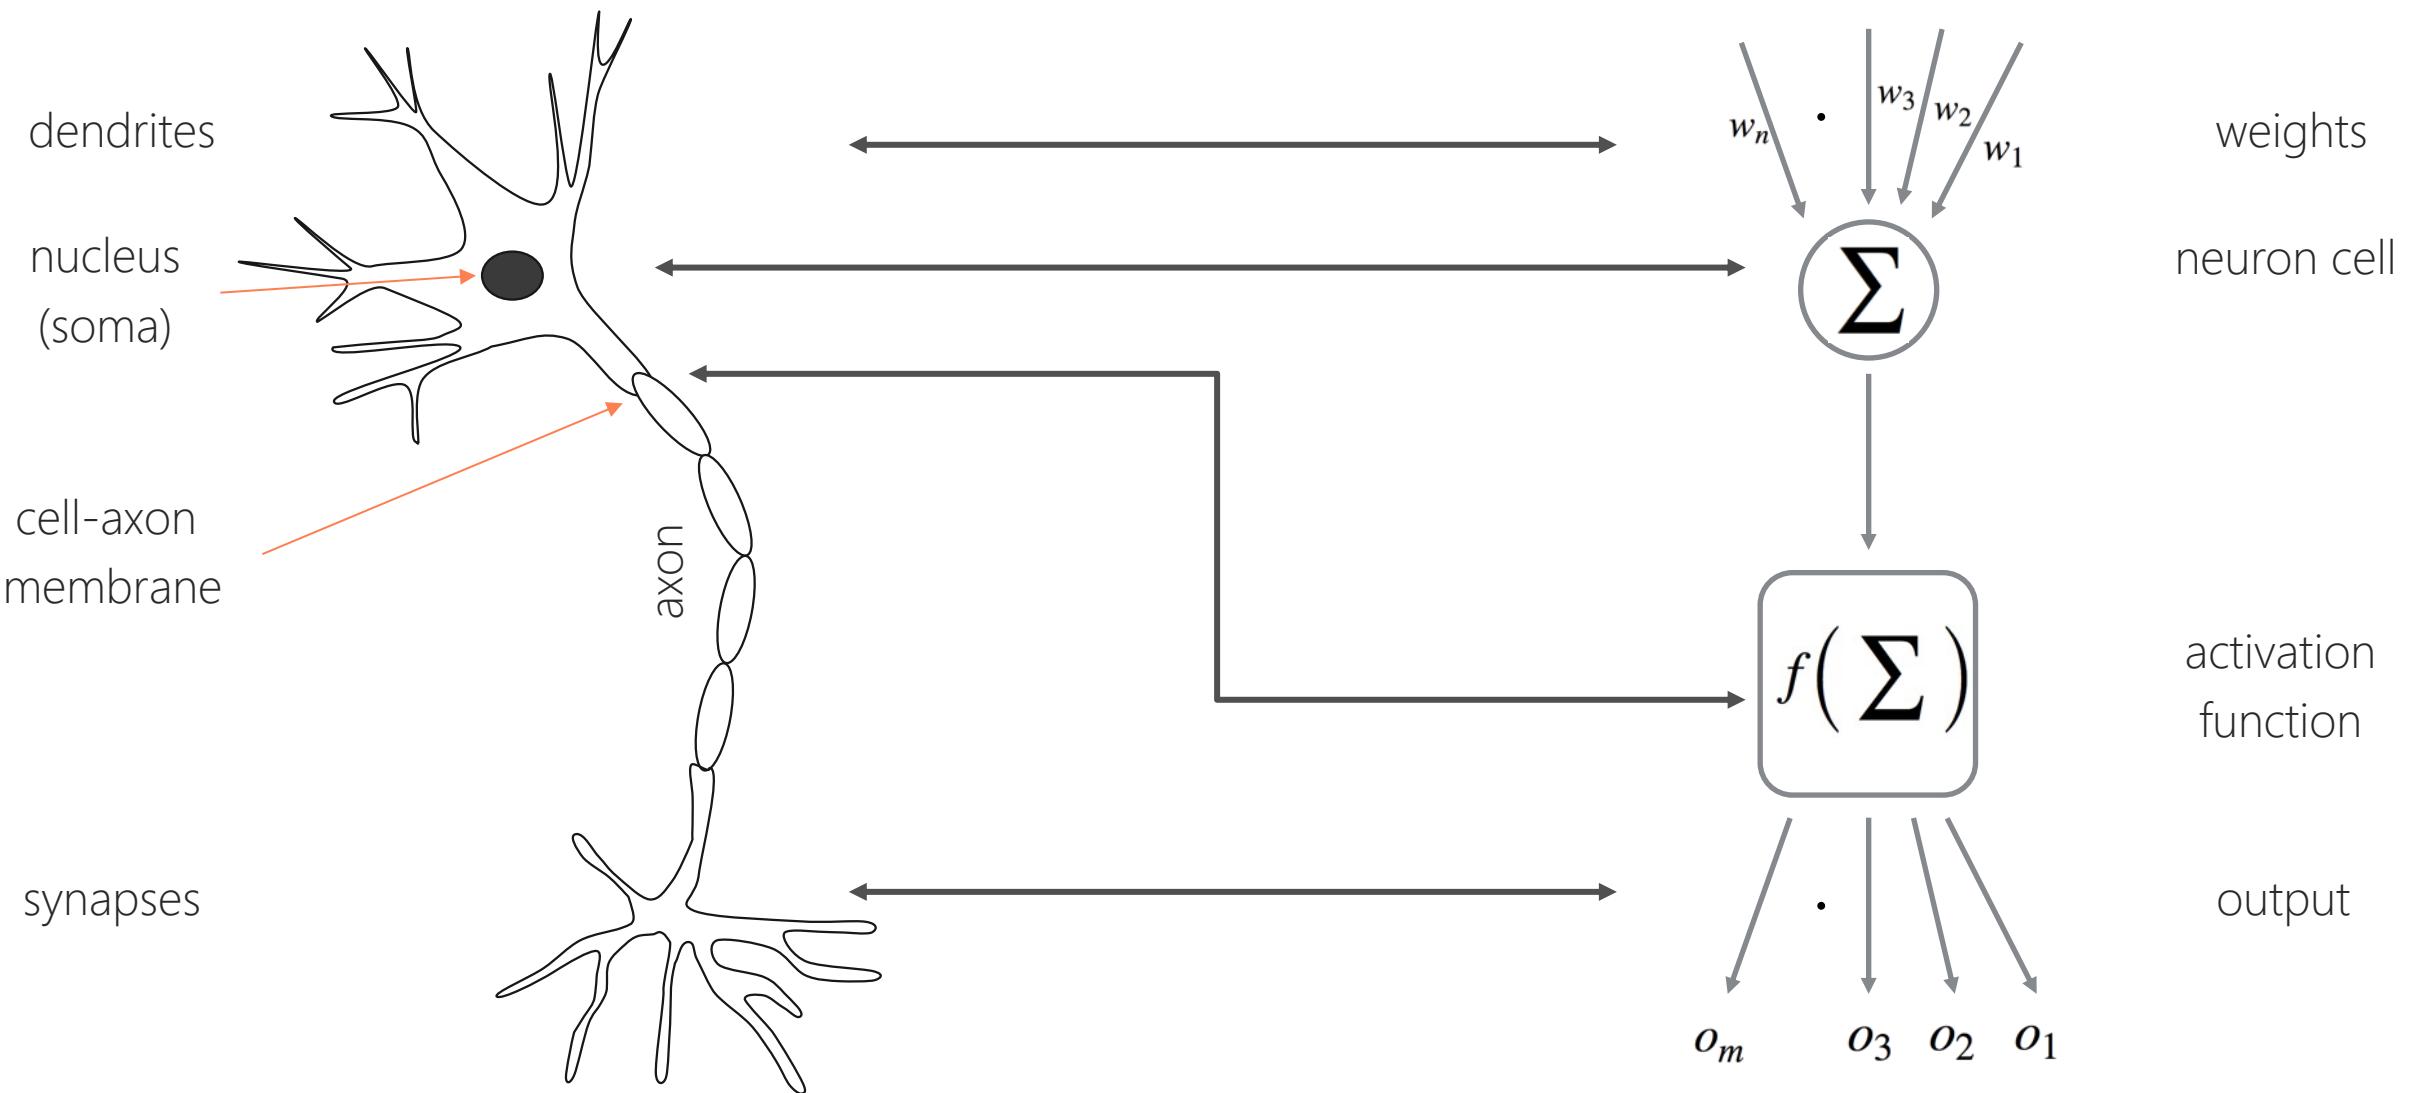
source: <http://bit.ly/2jjGINd>

Conceptual mathematical model

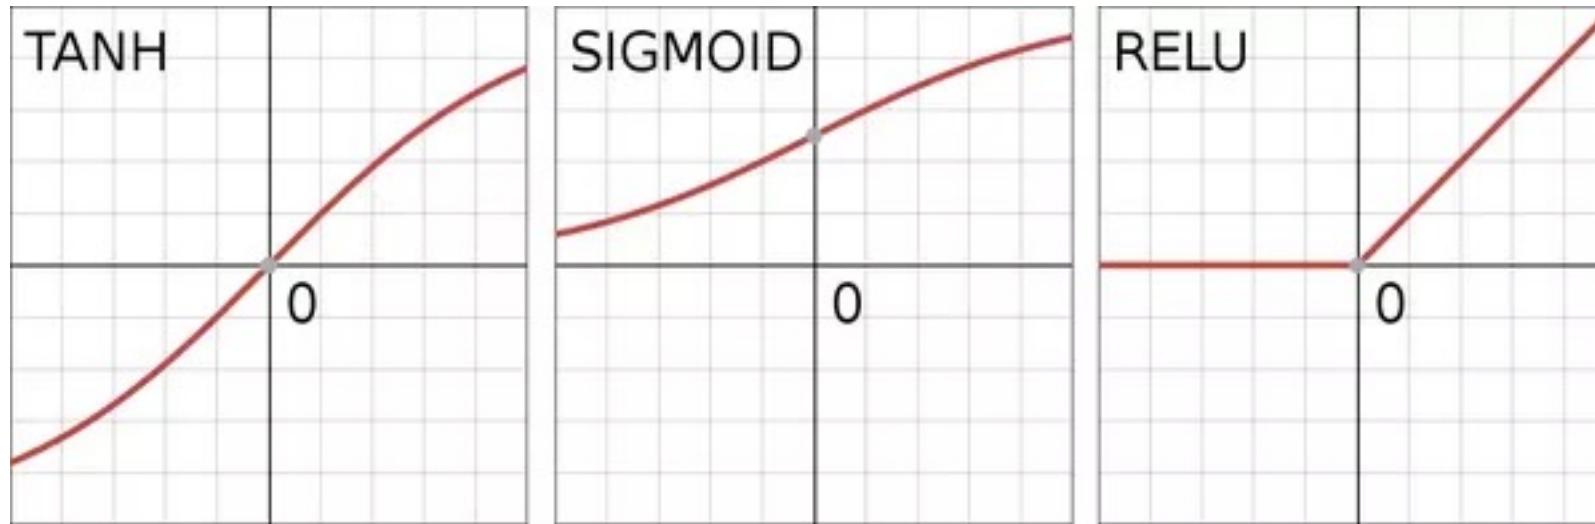
- Receives input from n sources
- Computes weighted sum
$$h_1 = x_1 w_1 + x_2 w_2 + \dots + x_n w_n$$
- Passes through an *activation function*
- Sends the signal to m succeeding neurons



Parallels

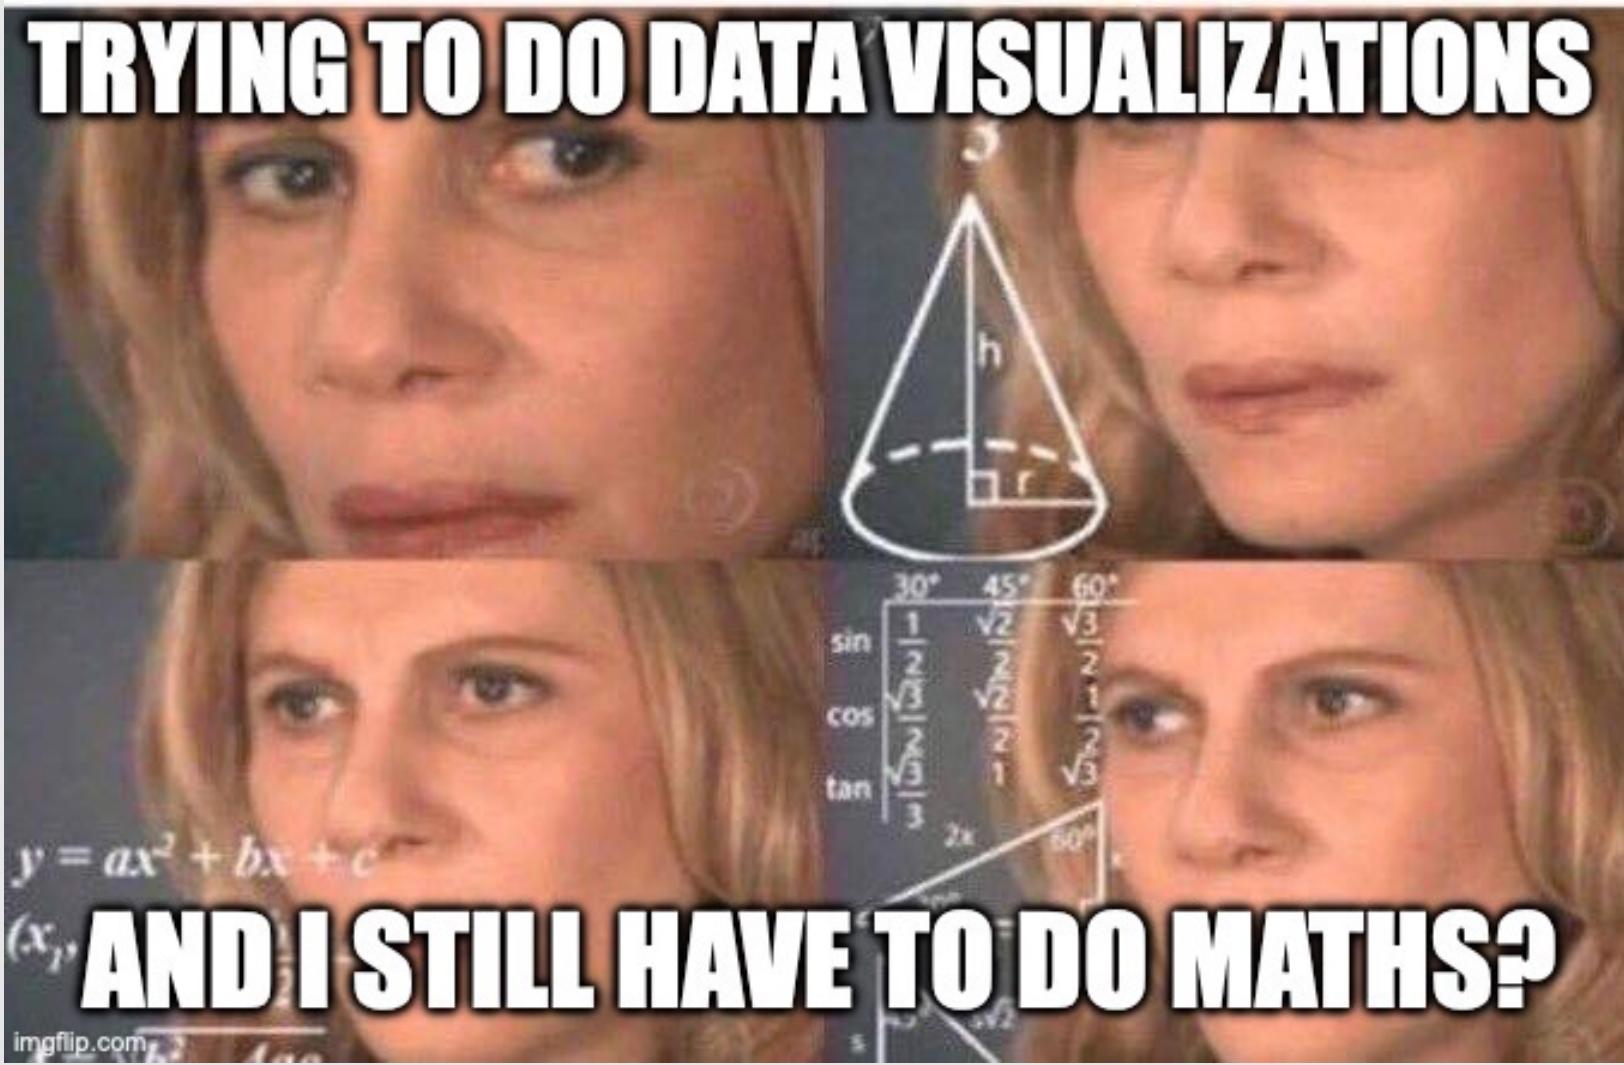


Activation functions

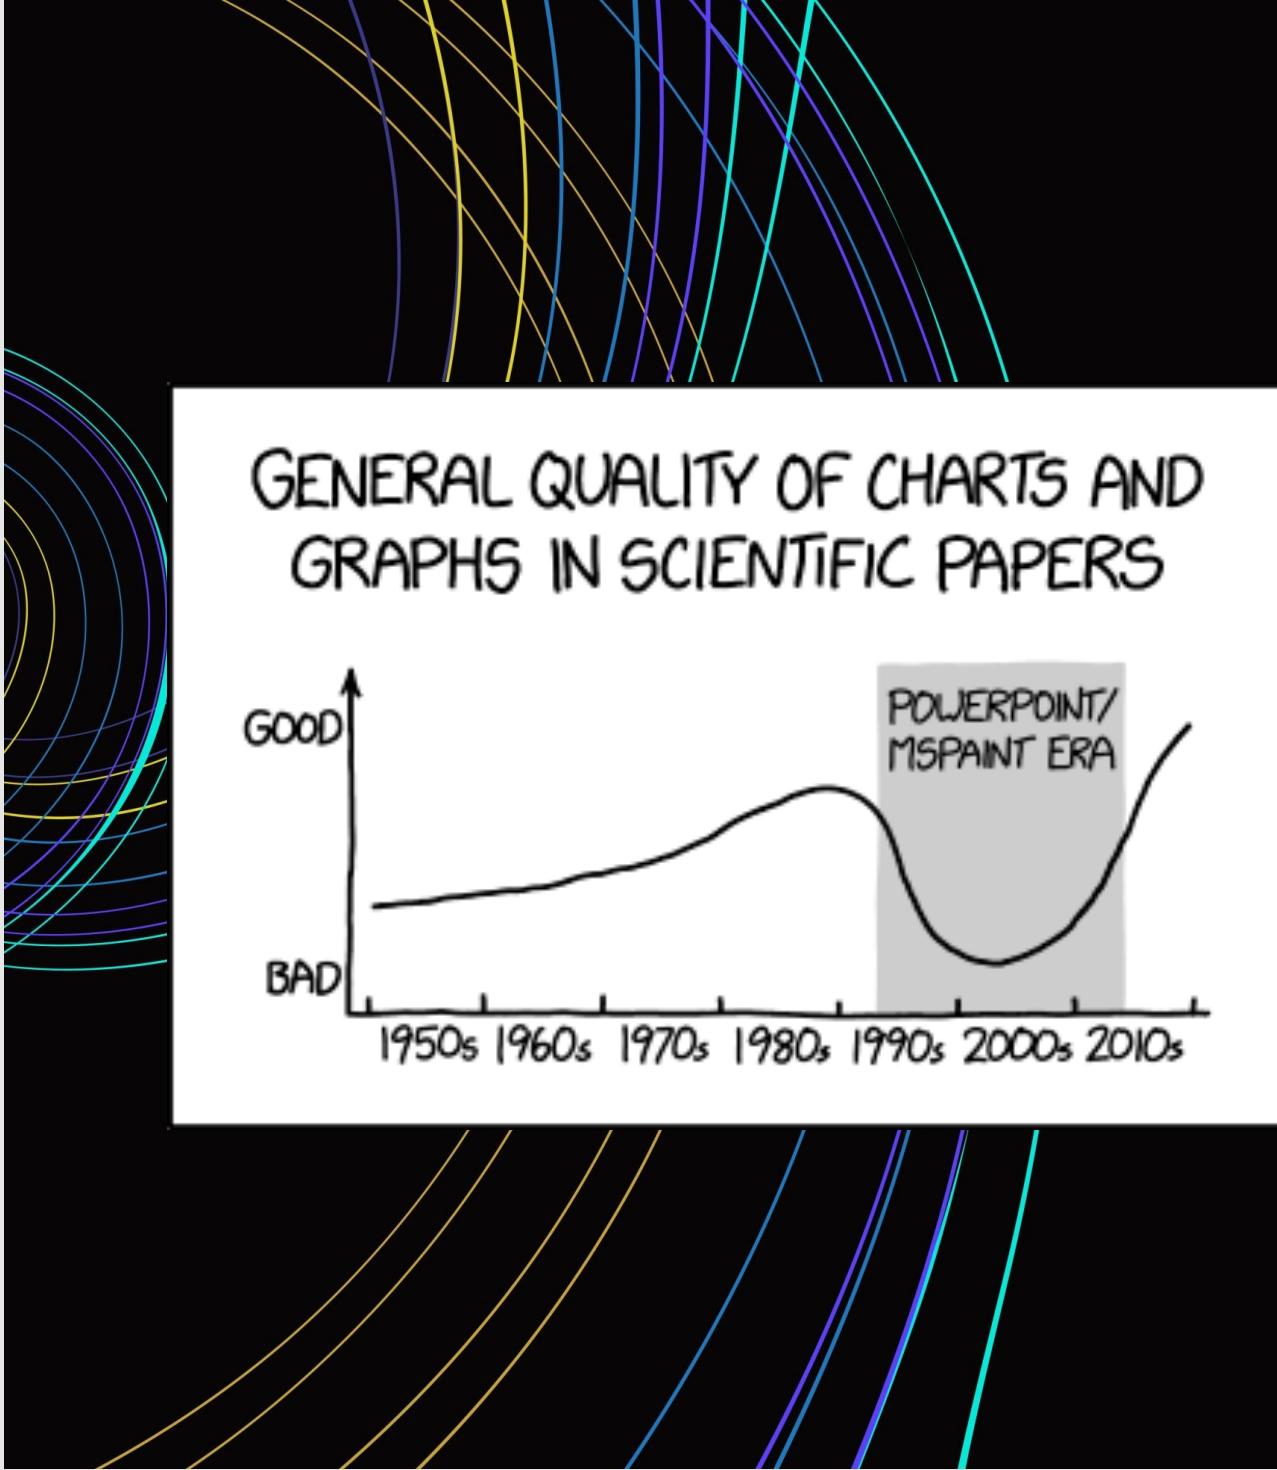


- Bias (threshold) activation function was proposed first
- Sigmoid and tanh introduce non-linearity with different codomains
- ReLu is the latest and greatest due to backprop (more later)

TRYING TO DO DATA VISUALIZATIONS



Don't worry, its
demo time

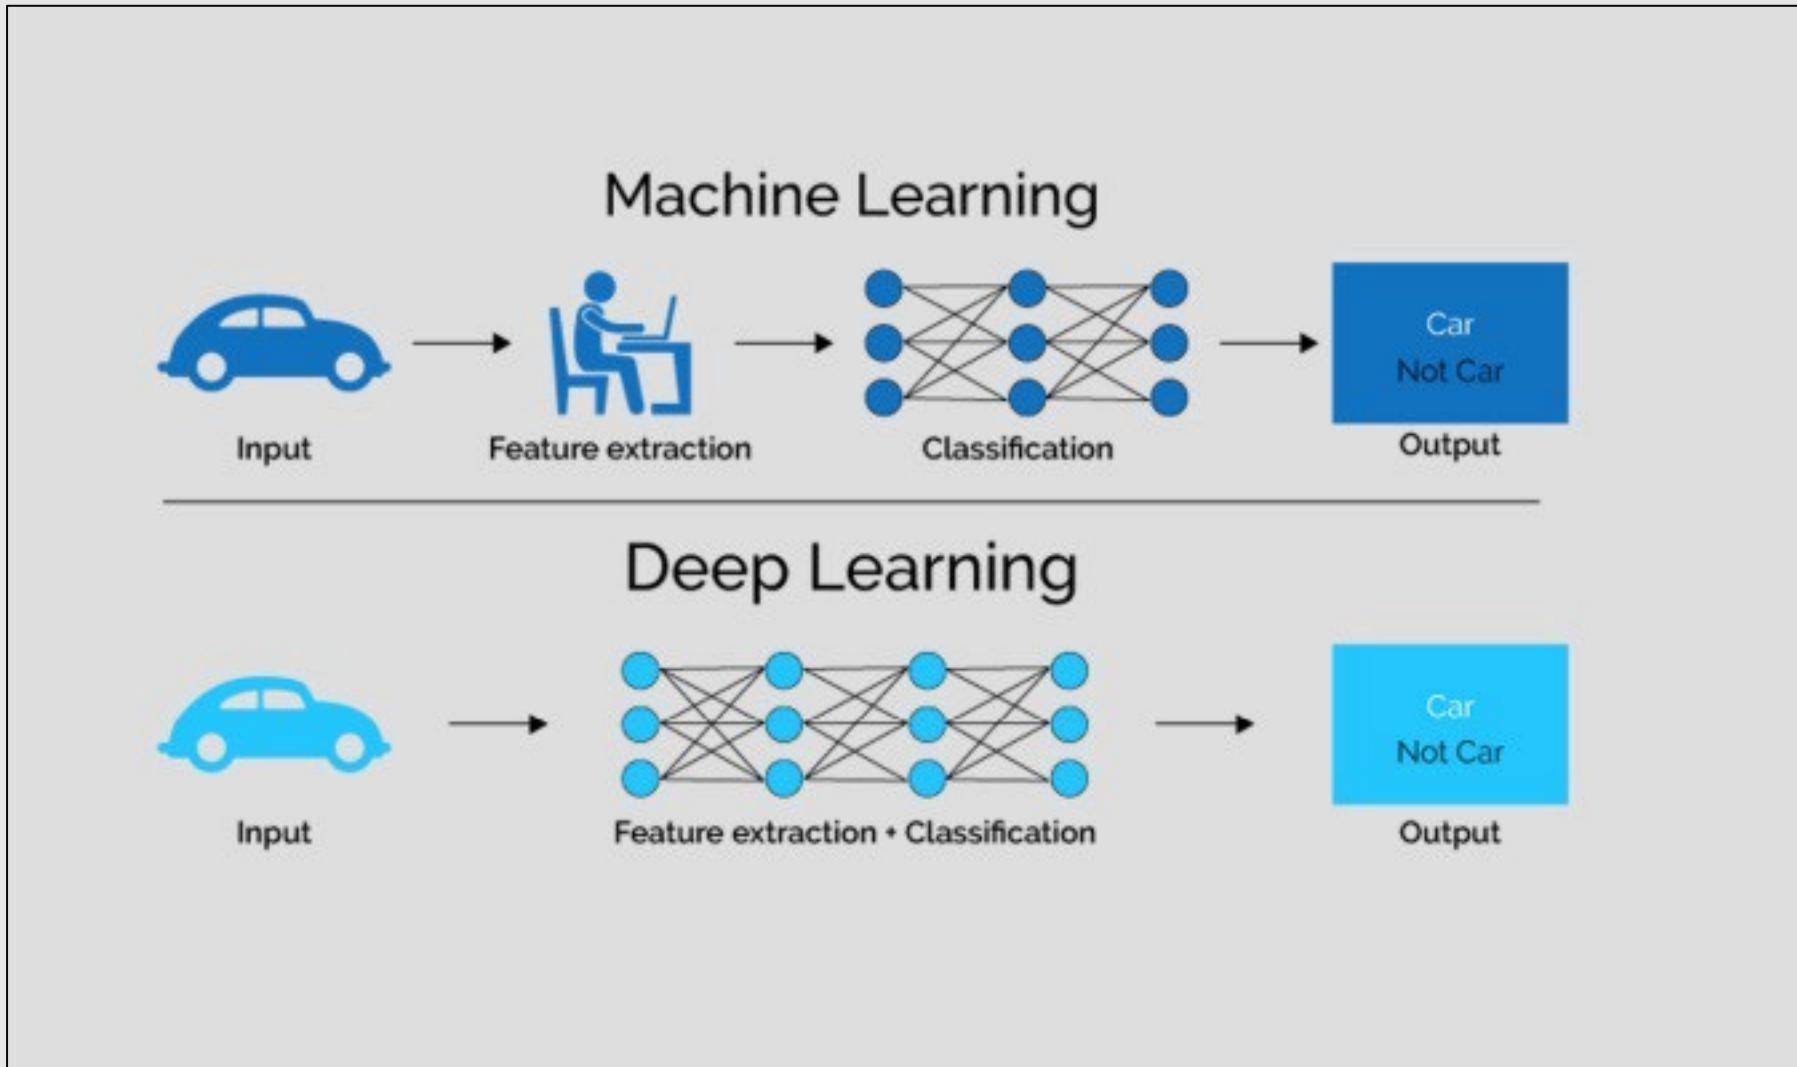


MNIST Keras Demo

Try it out at:

- Download the notebooks at Keras MNIST CNN notebooks: [CPU](#) | [GPU](#)
- The notebooks are for [Databricks Community Edition \(free\)](#) so work with your parents and teachers to get access

MNIST Keras Demo

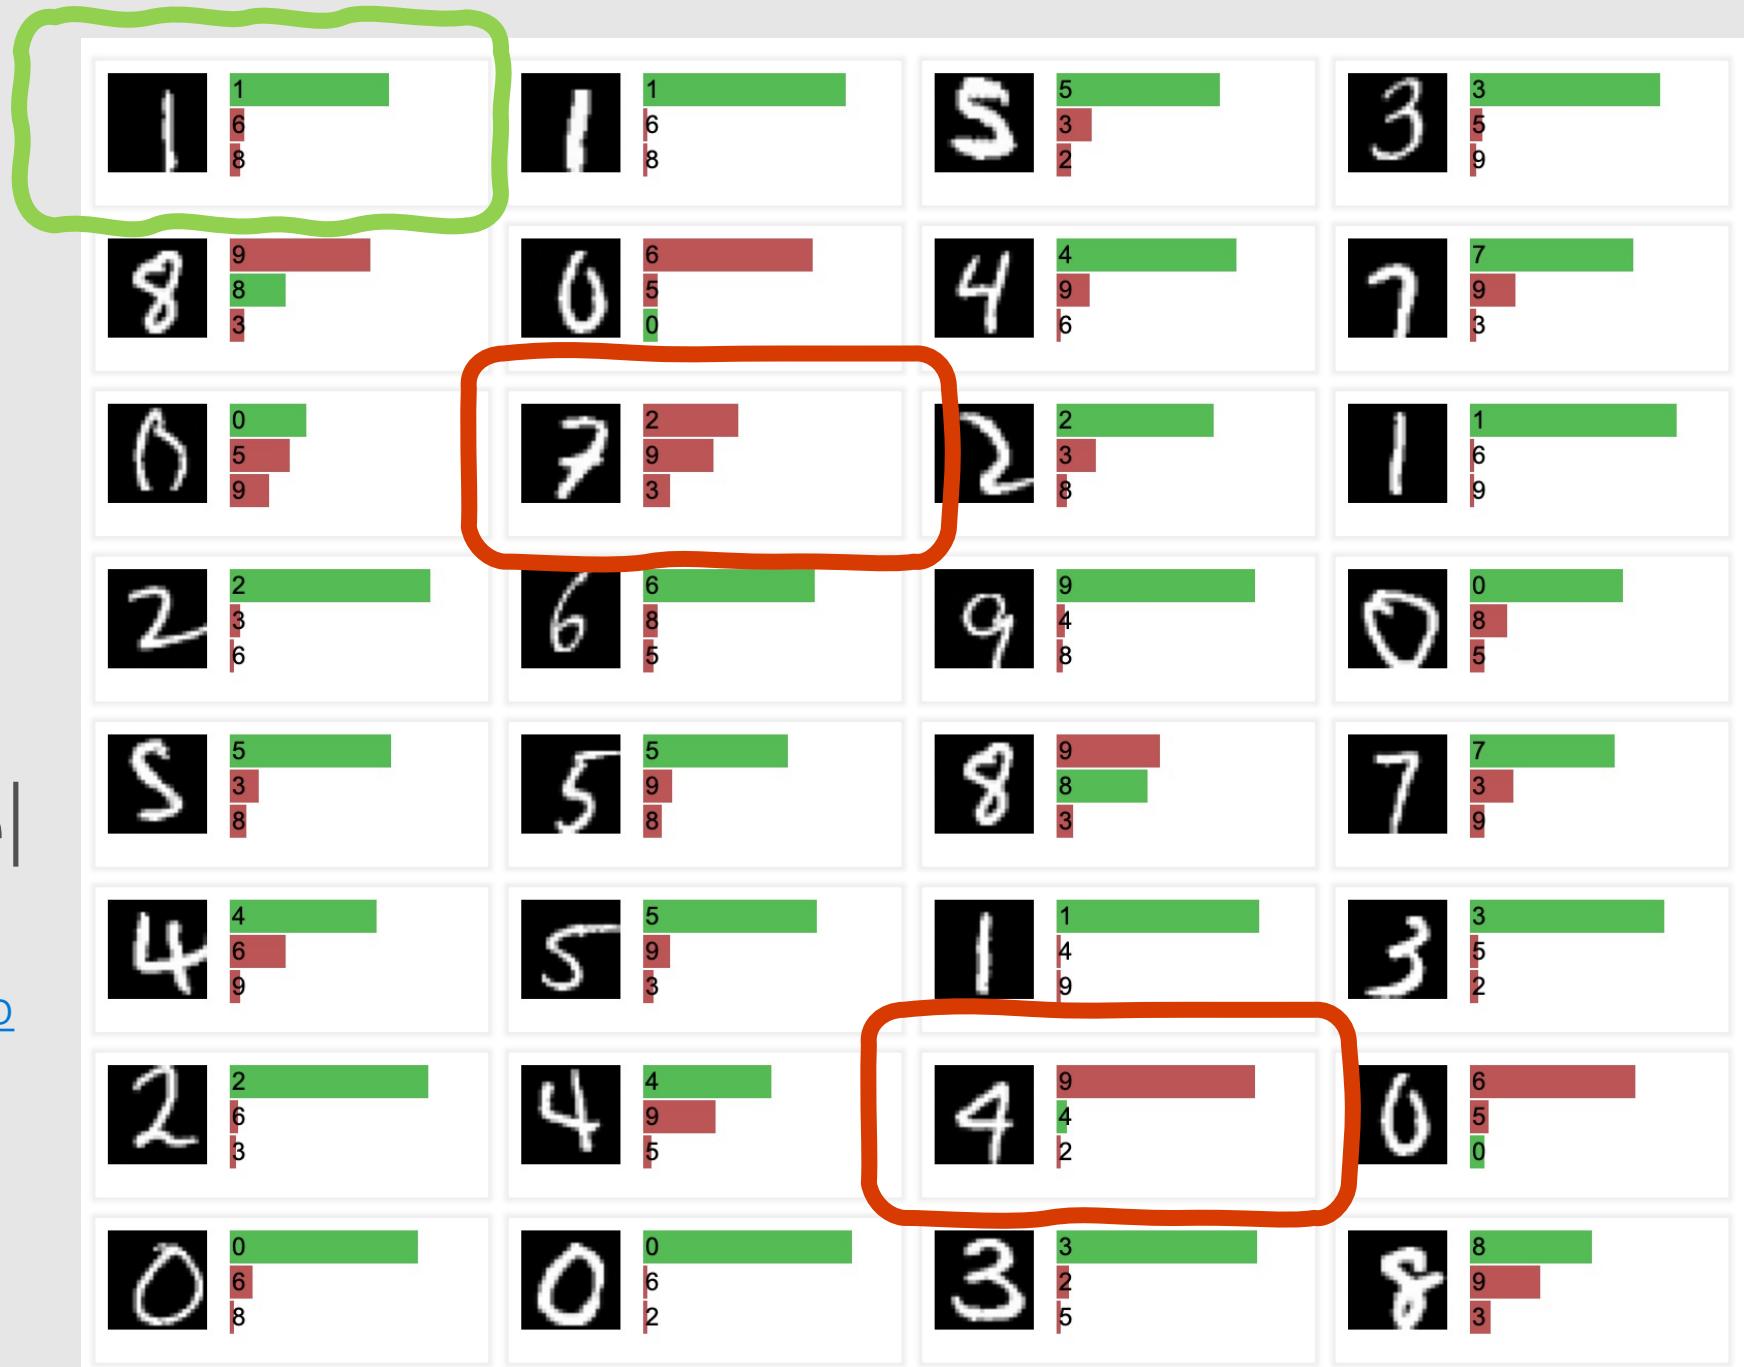


Source: [What is Training Data its types and why it is important?](#)

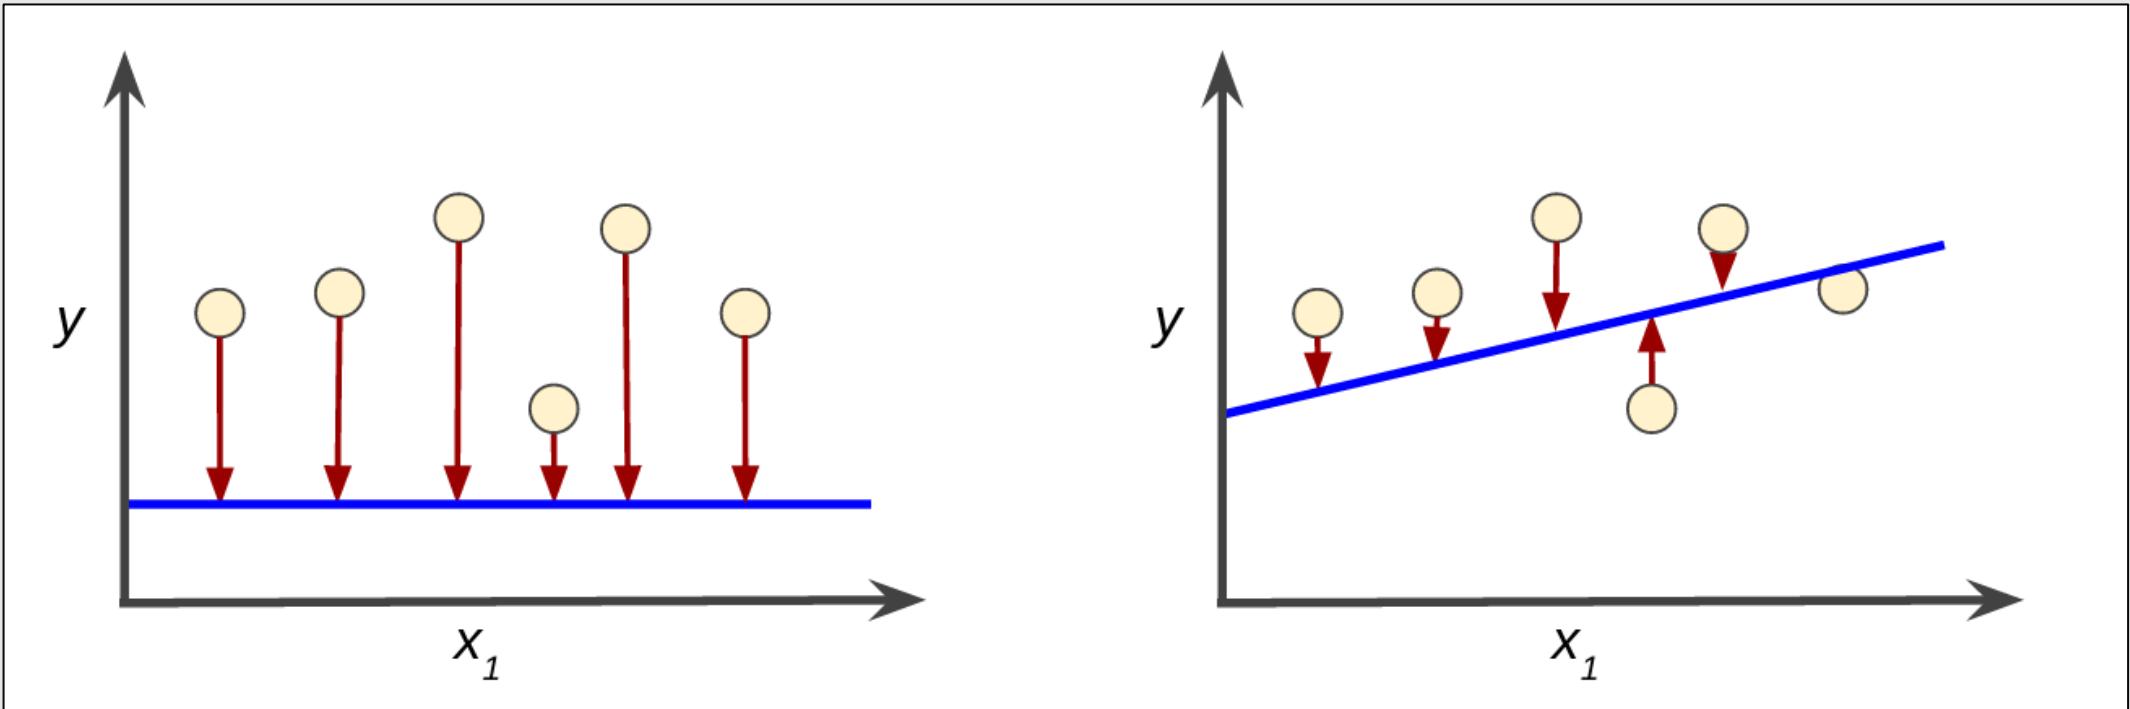
MNIST Keras Demo

More data so
we can better
train the model

Source: [ConvNetJS MNIST demo](#)



MNIST Keras Demo



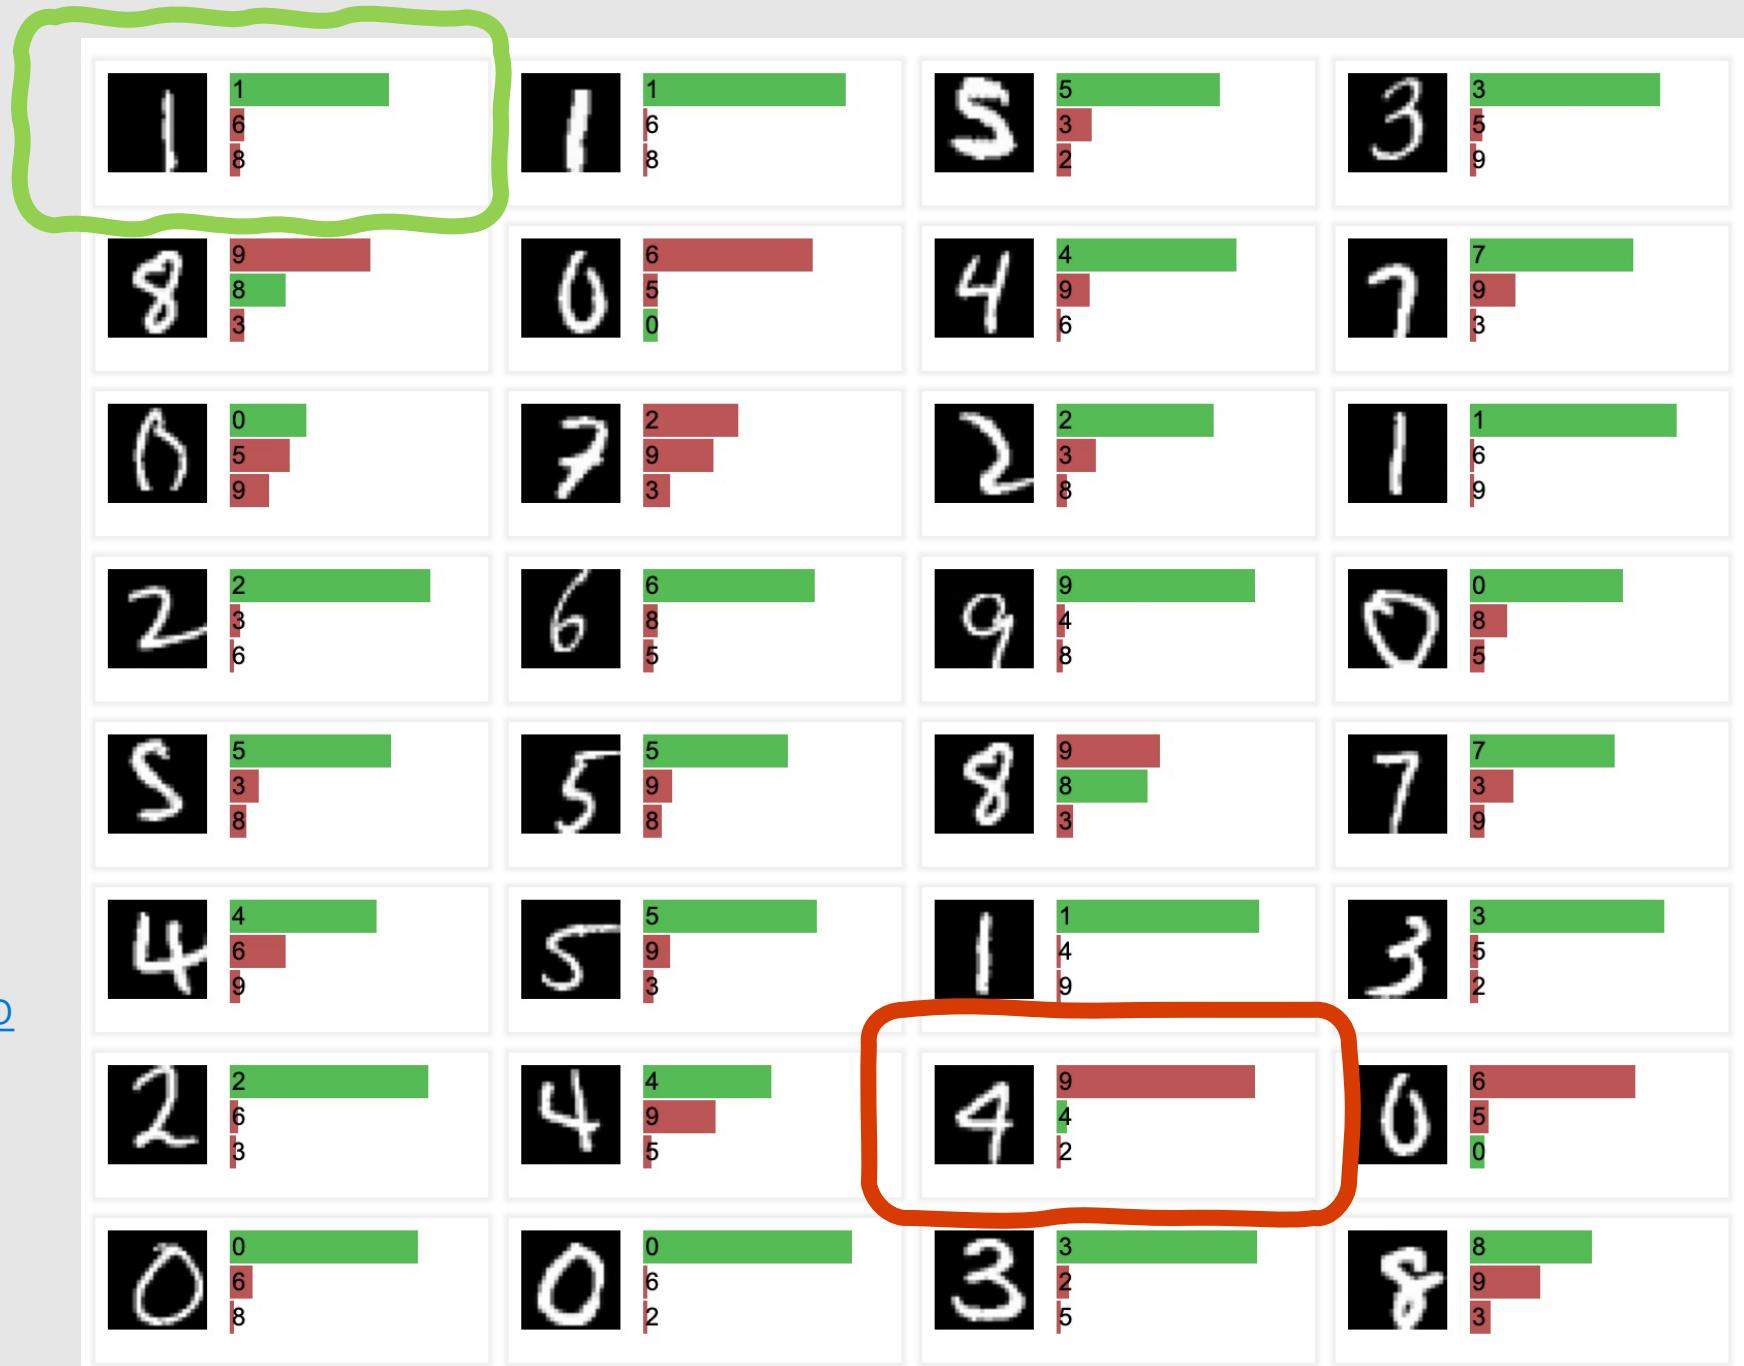
High loss in the left model; low loss in the right model.

Source: [Descending into ML: Training and Loss](#)

MNIST Keras Demo

Test Accuracy:
More green,
less red

Source: [ConvNetJS MNIST demo](#)



ConvNetJS MNIST demo



ConvNetJS CIFAR-10 demo

accuracy based on last 200 test images: 0.28191489361702127



Deep Visualization Toolbox



Source: <https://youtu.be/AgkfIQ4IGaM?t=118>

Object Detection

Image AI: Video Object Detection, Tracking and Analysis utilizing YOLOv3 model



Source: [Coffee Object Detection Demo](#)

DALL·E 2

is a new AI system that can create realistic images and art from a description in natural language.

*An astronaut
riding a horse in a
photorealistic style*

<https://openai.com/dall-e-2/>



Want to know more about DALL·E 2?

How DALL·E 2 actually works?

- Link text and visual semantics via [CLIP](#)
 - *The insight that makes CLIP so powerful is the bright idea to train models to not just identify which category (of a pre-defined list of options) an image belongs to, but to identify the caption of each image from a list of random captions.¹*
- Generating Images from Visual Semantics by doing the inverse with [GLIDE](#) via a diffusion model

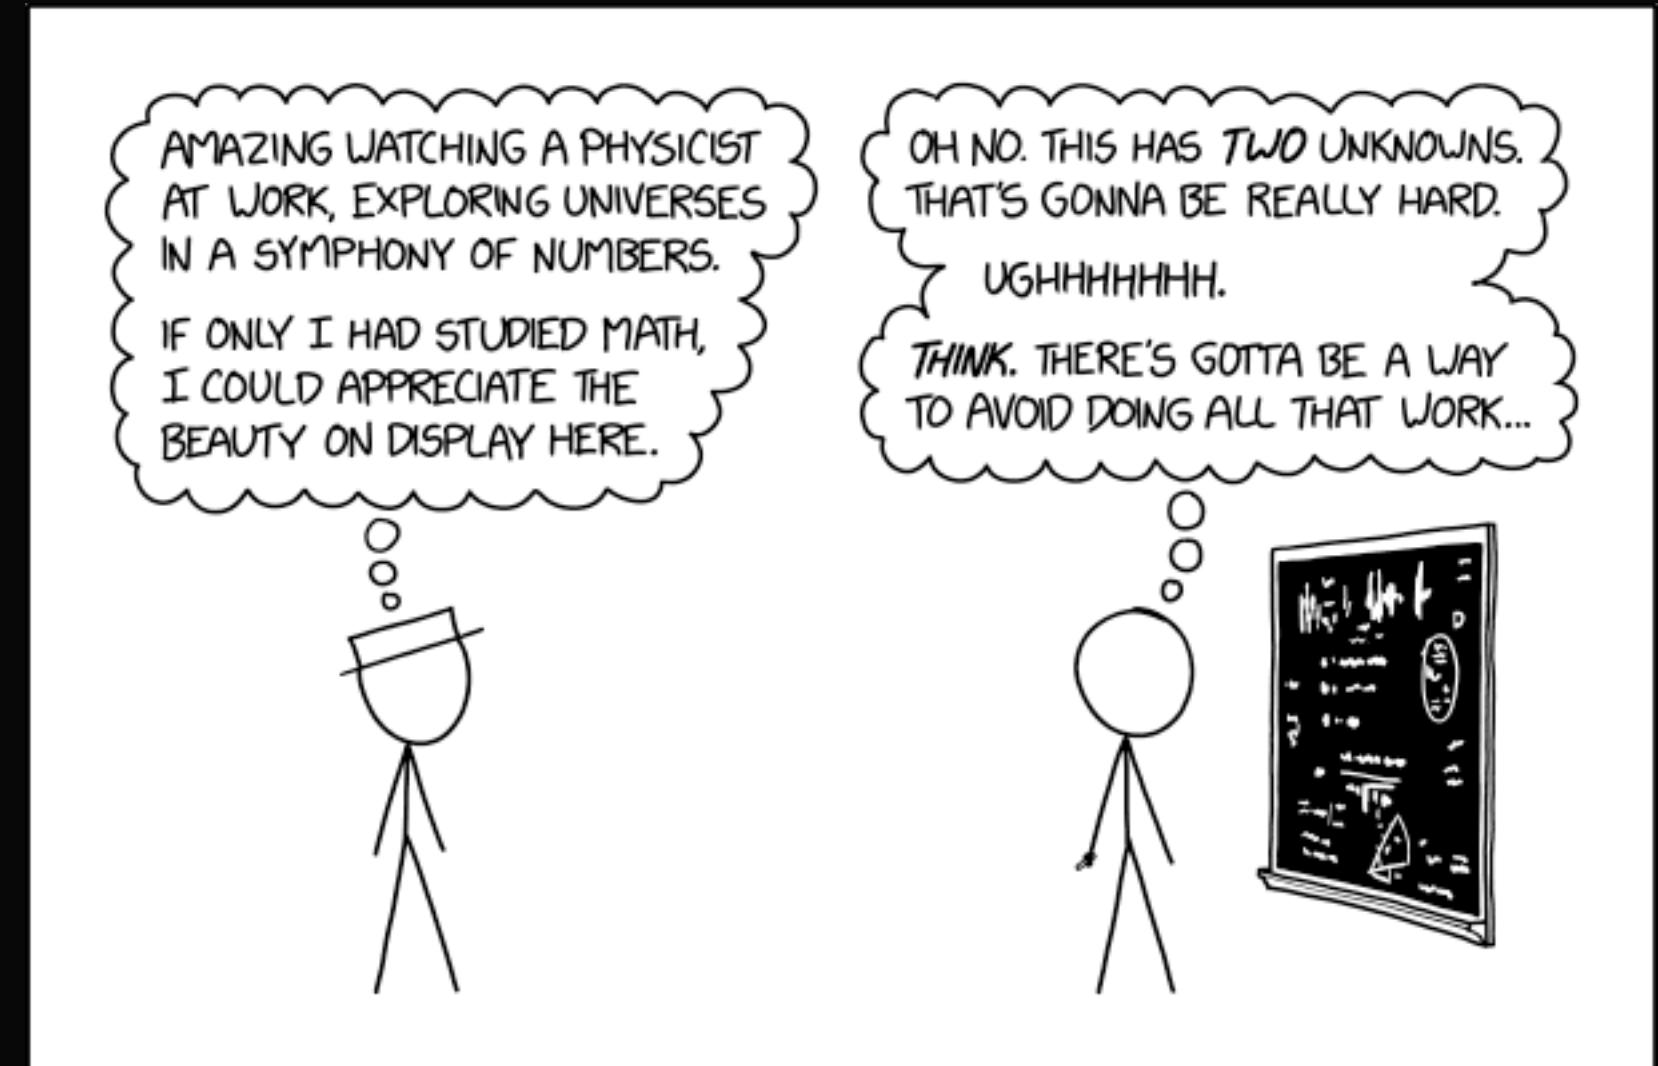
Want to know more about DALL·E 2?

What are diffusion models?

- *Diffusion Models learn to generate data by reversing a gradual noising process.*
- Think of this as unscrambling a Rubik's cube after scrambling it one step at at time (after forgetting about the steps) or
- *More practically, you could train a neural network to go from a state of disorder to a state of less disorder.*

¹ [DALL·E 2.0, Explained](#)

Want to do
more
math?



References

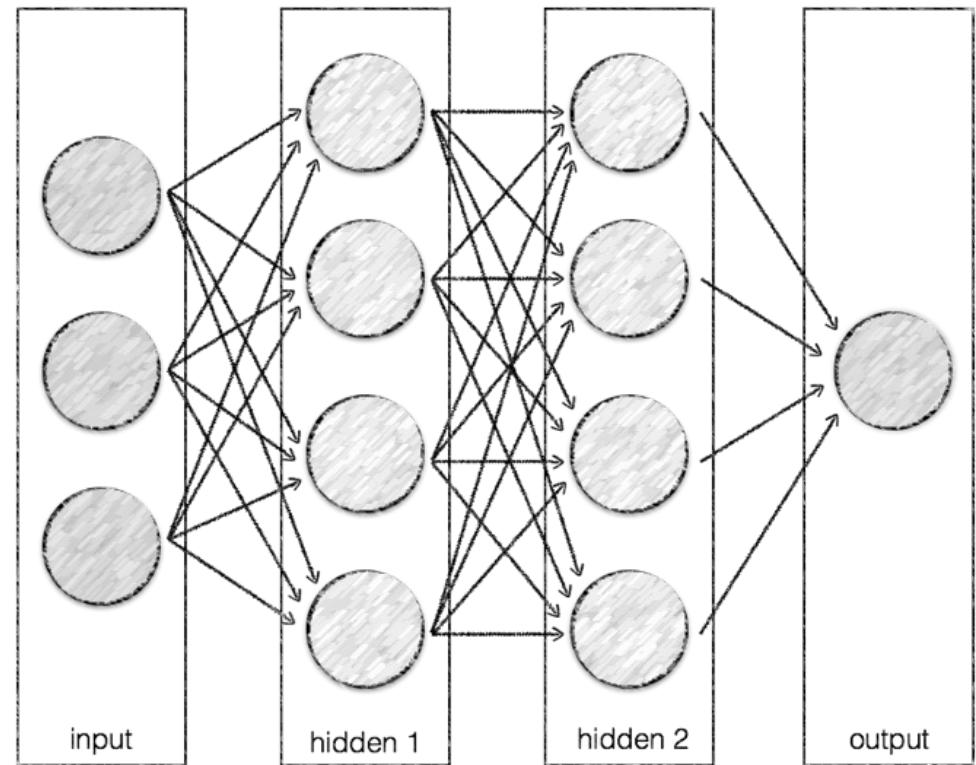
These slides: [Data Visualization – Northstar Middle School](#)

Keras MNIST Notebooks: [GPU](#) | [CPU](#)

Neural networks introduction

Artificial Neural Network

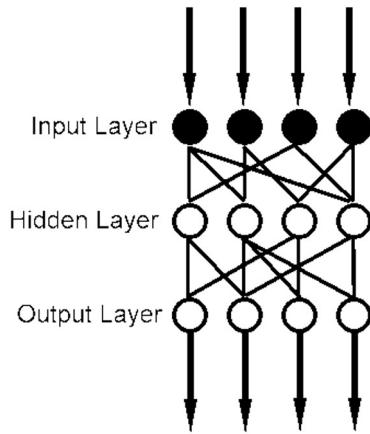
- Organized into layers of neurons
- Typically 3 or more: input, hidden and output
- A black-box model



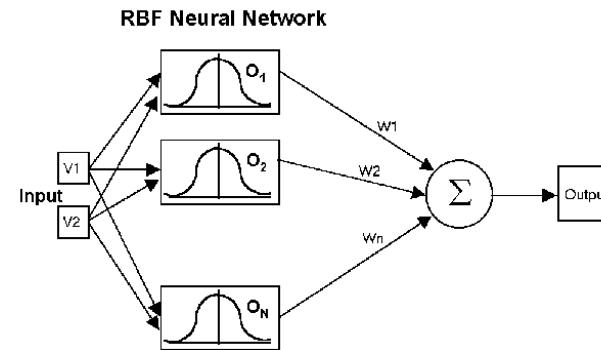
Types of neural networks

Some examples

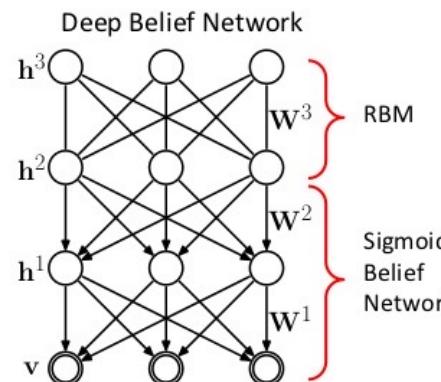
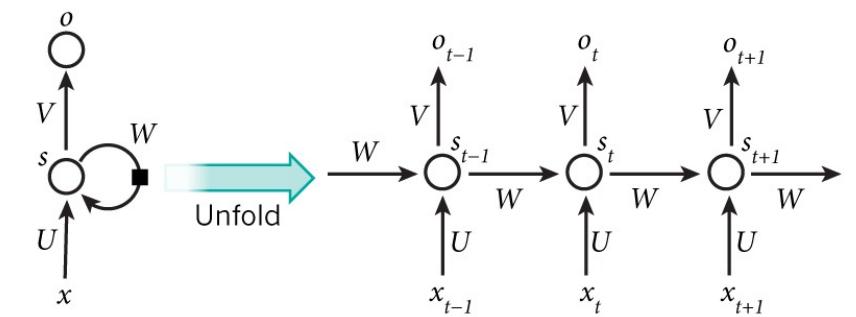
Single- or multi-layer feedforward networks



Radial Basis Function



Recurrent Neural Networks

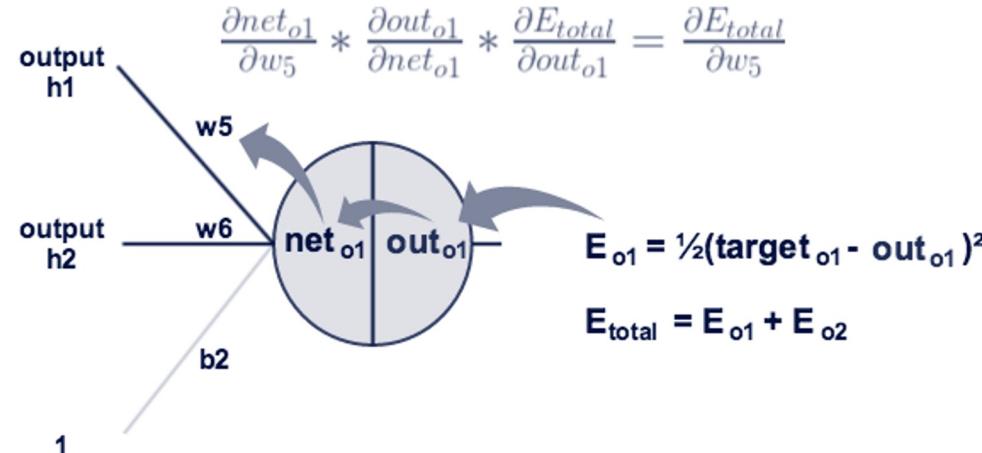


Deep belief networks

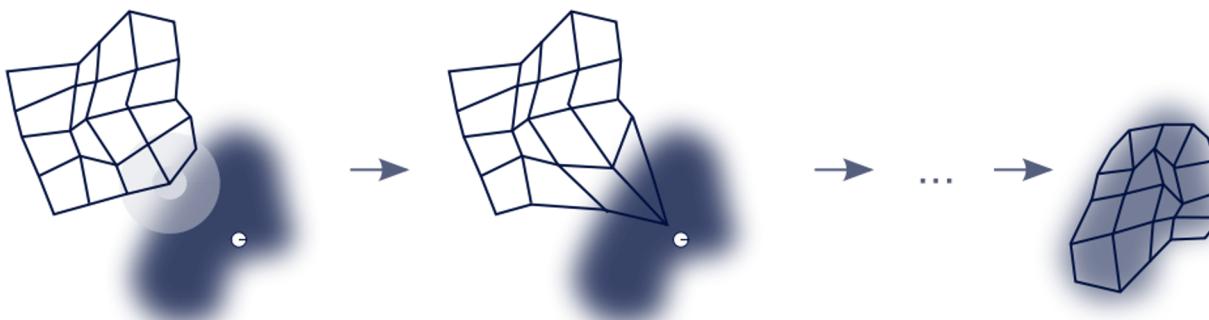
Learning

Some examples

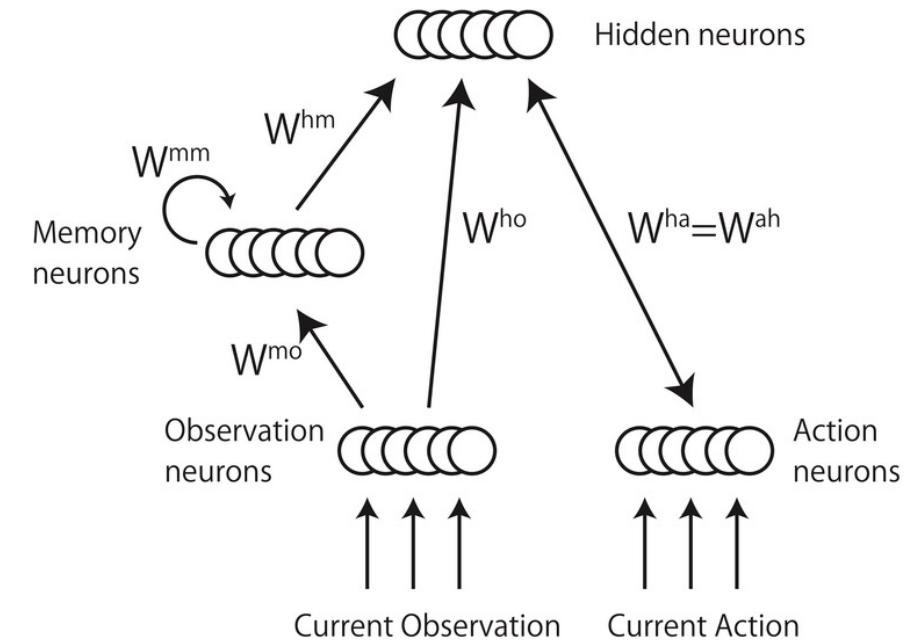
Backpropagation



Self-organizing maps (competitive learning)



Memory based learning



Introduction to tensors

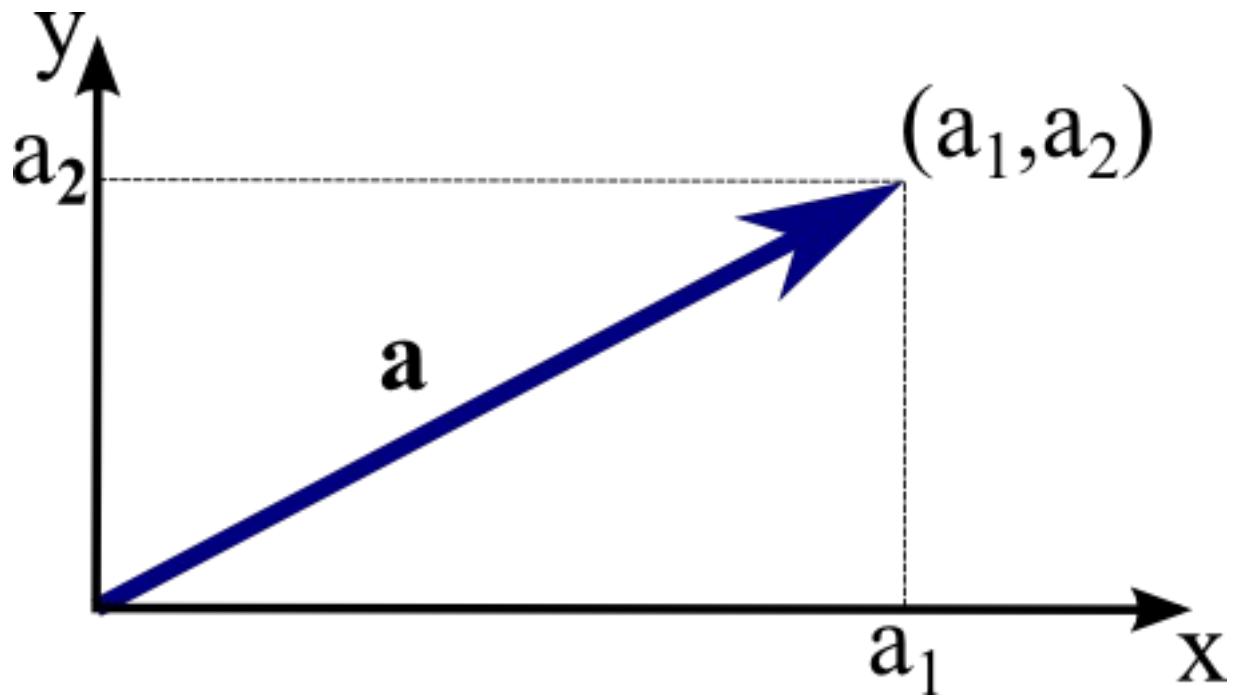
Scalars, vectors, matrices and tensors

Scalar

2.01

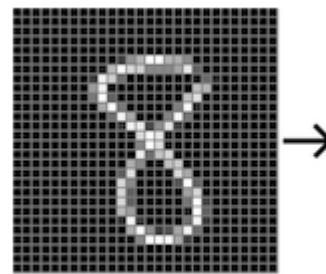
Scalars, vectors, matrices and tensors

Vector



Scalars, vectors, matrices and tensors

Matrix



28 x 28
784 pixels

Scalars, vectors, matrices and tensors

Tensor (3D)



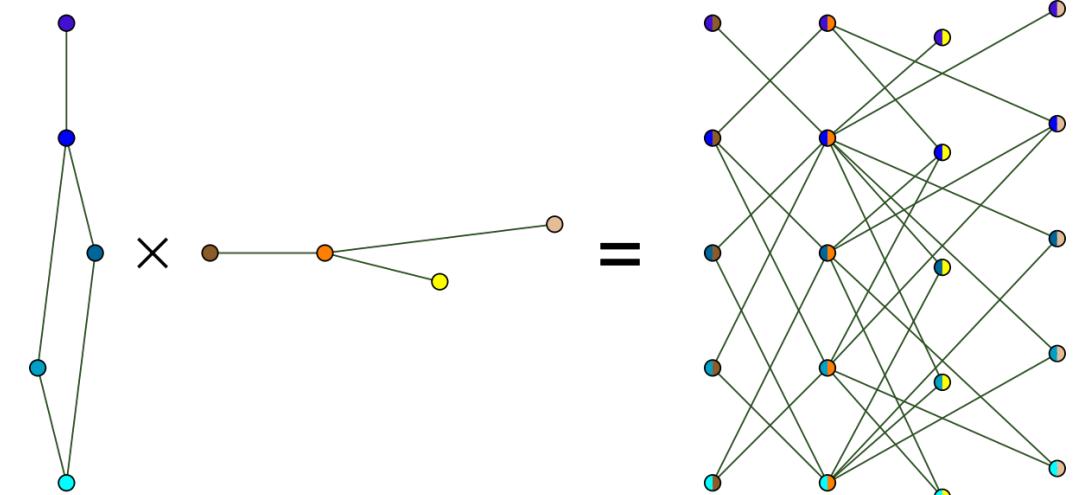
Scalars, vectors, matrices and tensors

Tensor (4D)

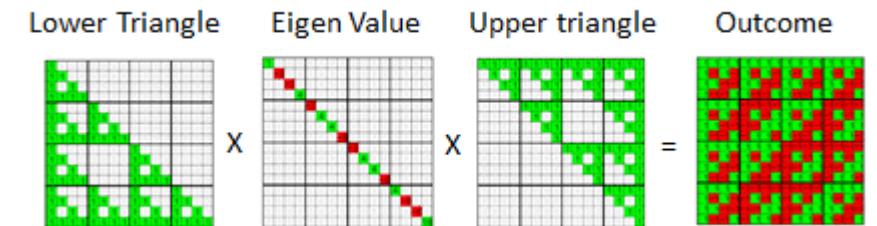


Tensor operations

- Scaling
- Addition and subtraction
- Multiplication
- Contraction
- Tensor product
- Reshaping
- Transposition
- Eigenvectors and eigenvalues
- Cholesky decomposition
- ...



Source: <http://bit.ly/2F7svG0>



Source: <http://bit.ly/2o8dyvy>

Concept

- t_1 is a 2×3 matrix
- t_2 is a 3×2 matrix
- The op_1 is a multiplication operator
- The above can be represented in a matrix form

$$t_1 = \begin{bmatrix} a & b & c \\ d & e & f \end{bmatrix}$$
$$t_2 = \begin{bmatrix} u & v \\ w & x \\ y & z \end{bmatrix}$$

A diagram illustrating the multiplication of two matrices. On the left, there is a matrix t_1 with elements a, b, c in the first row and d, e, f in the second row. On the right, there is a matrix t_2 with elements u, v in the first row, w, x in the second row, and y, z in the third row. Two arrows point from these matrices to a blue-outlined circle containing the label op_1 .

$$out = t_1 \underset{op_1}{*} t_2 = \begin{bmatrix} au + bw + dy & av + bx + cz \\ du + ew + fy & dv + ex + fz \end{bmatrix}$$

Tensor product

- t_1 is a 2×3 matrix
- t_2 is a 3×2 matrix
- The op_1 is a tensor product operator
- The above can be represented in a matrix form

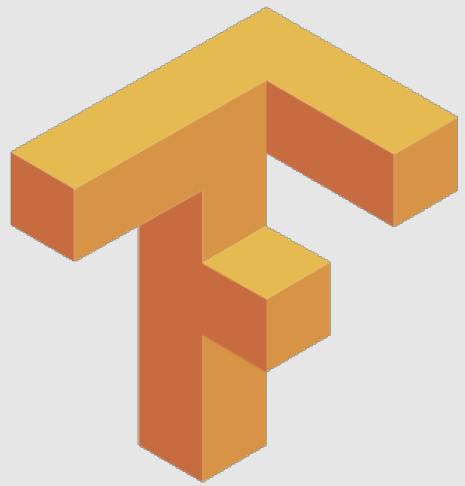
$$out = t_1 \underset{op_1}{\otimes} t_2 = \begin{bmatrix} at_2 & bt_2 & ct_2 \\ dt_2 & et_2 & ft_2 \end{bmatrix}$$

$$t_1 = \begin{bmatrix} a & b & c \\ d & e & f \end{bmatrix} \quad t_2 = \begin{bmatrix} u & v \\ w & x \\ y & z \end{bmatrix}$$

A diagram illustrating the tensor product. Two matrices, t_1 and t_2 , are shown on the left. Arrows point from each matrix to a central blue circle containing the label op_1 .

$$\begin{bmatrix} au & av & bu & bv & cu & cv \\ aw & ax & bw & bx & cw & cx \\ ay & az & by & bz & cy & cz \\ du & dv & eu & bv & fu & fv \\ dw & dx & ew & bx & fw & fx \\ dy & dz & qy & bz & fy & fz \end{bmatrix}$$

Neural network as a graph

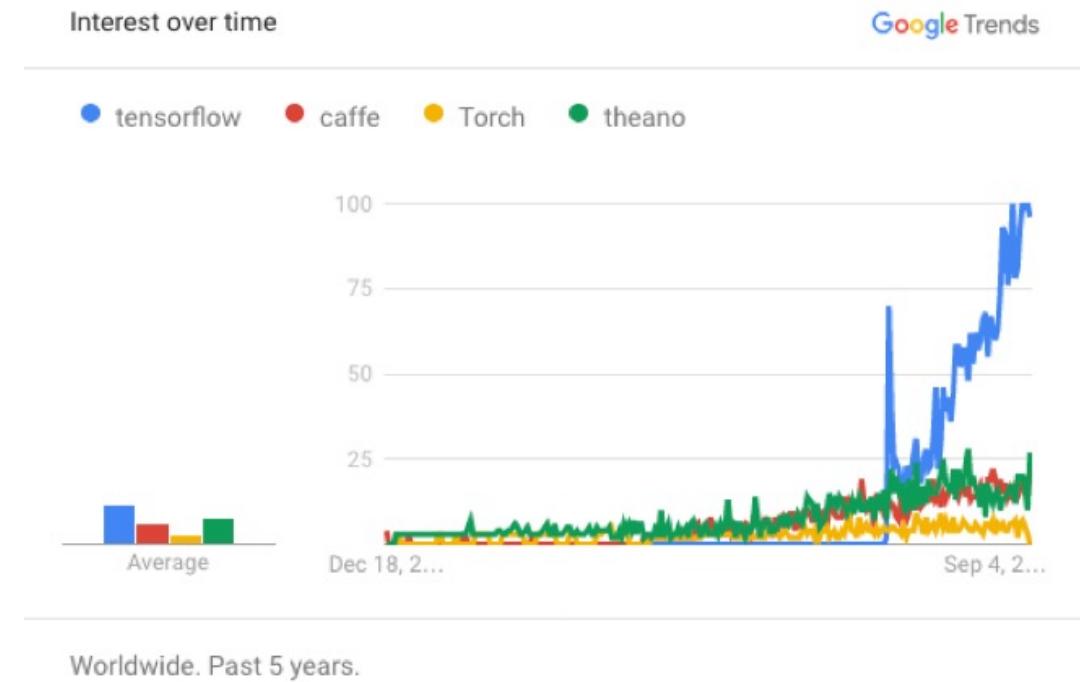


TensorFlow

Technology from Google

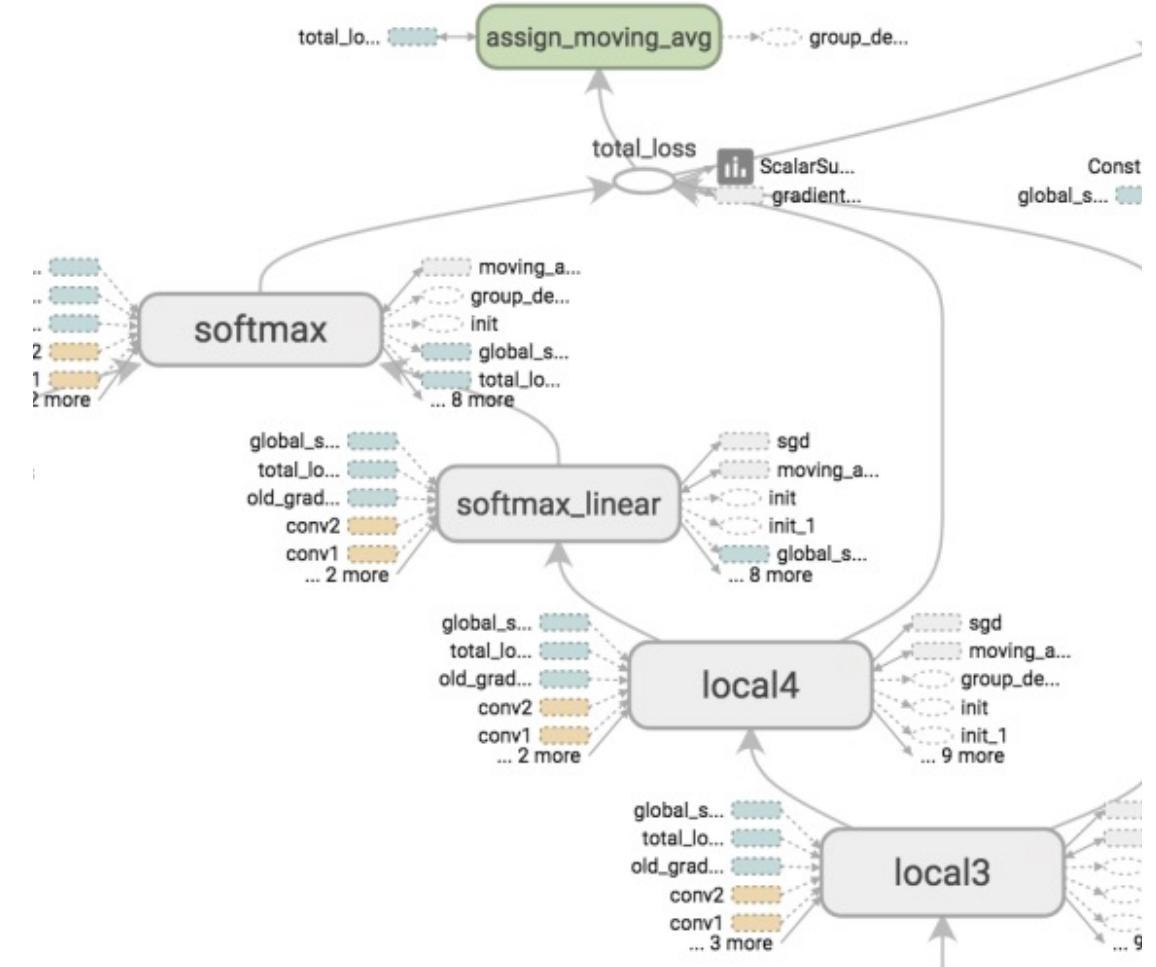
What is TensorFlow

- Open source software library for numerical computation using data flow graphs
- Based loosely on neural networks
- Built in C++ with a Python interface
- Quickly gained high popularity
- Supports GPU computations



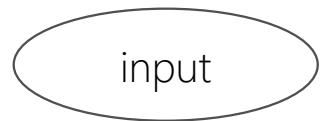
Data Flow Graph

- Nodes represent mathematical operations
- Edges are multidimensional arrays (tensors)
- Thus, TensorFlow is a flow of multidimensional arrays (tensors) between the mathematical operations (nodes)
- i.e. flow of *tensors*



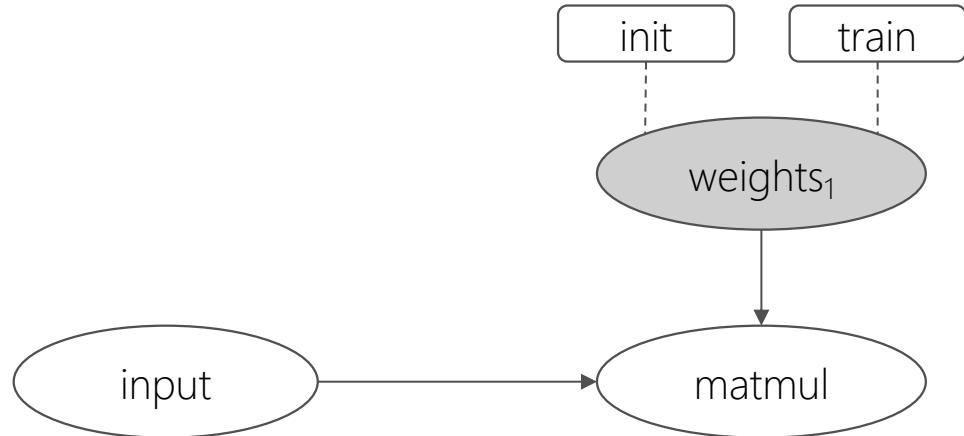
Step by step DAG: input

$$y = x$$



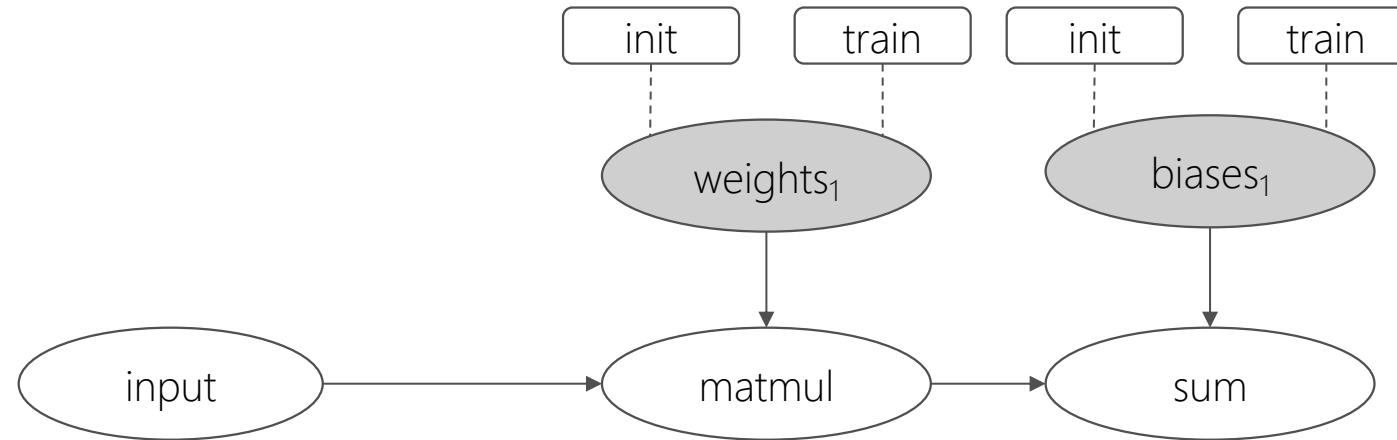
Step by step DAG: input

$$y = wx$$



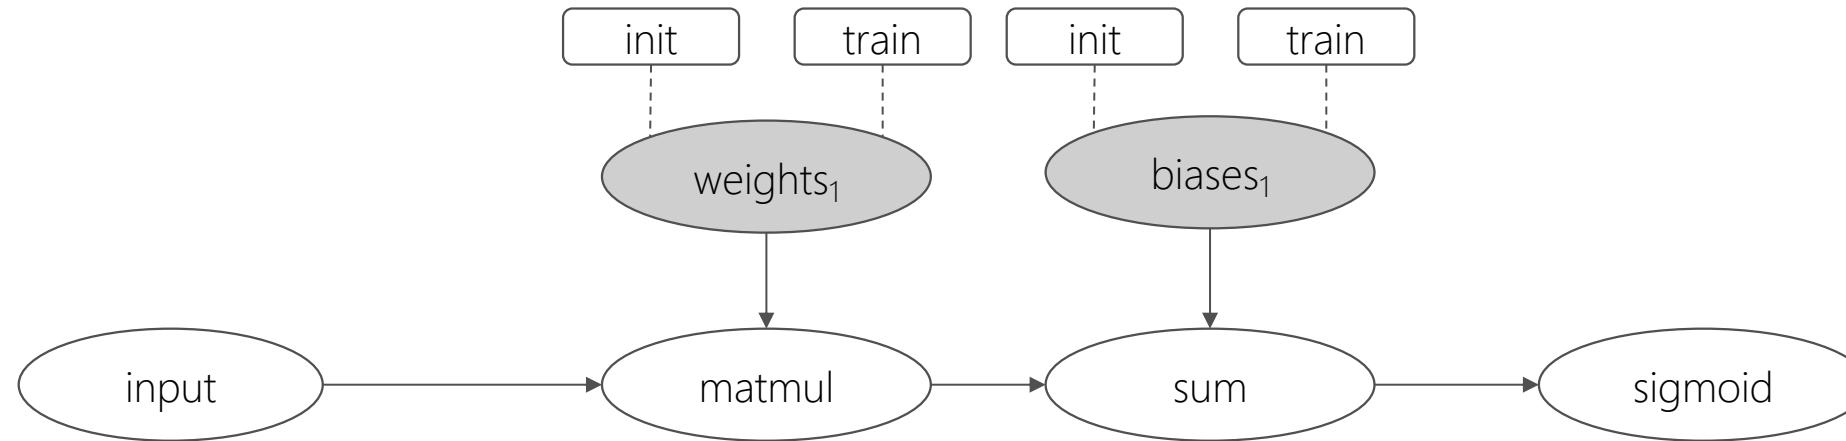
Step by step DAG: bias addition

$$y = wx + b$$



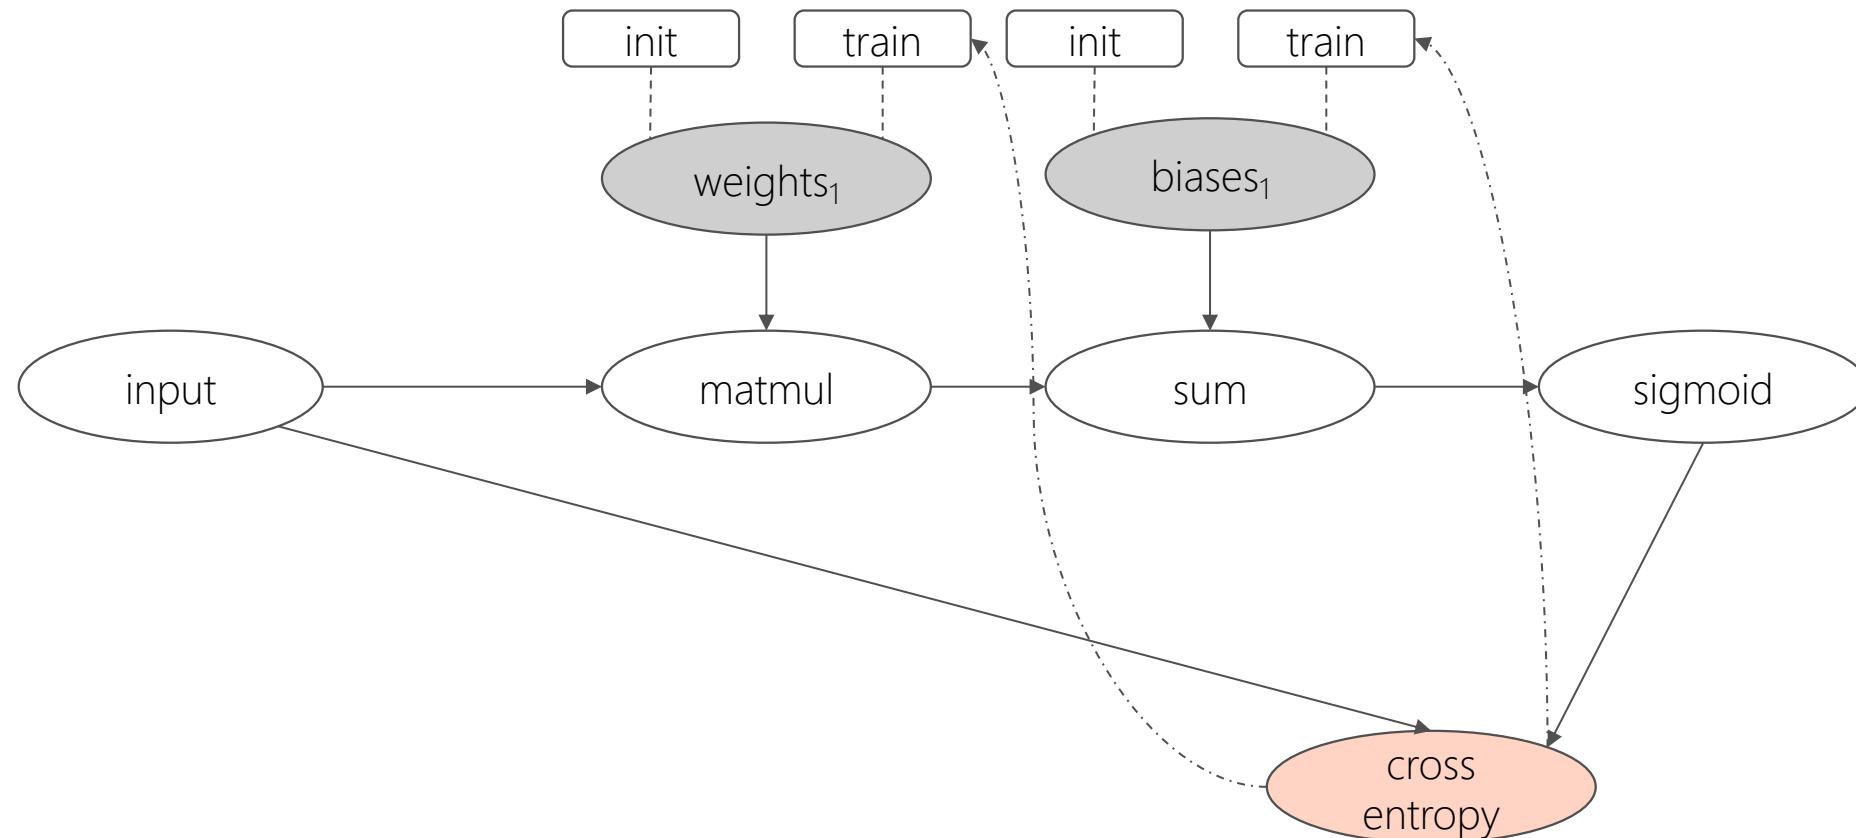
Step by step DAG: activation function

$$y = \frac{1}{1 + e^{-wx-b}}$$



Step by step DAG: cross entropy

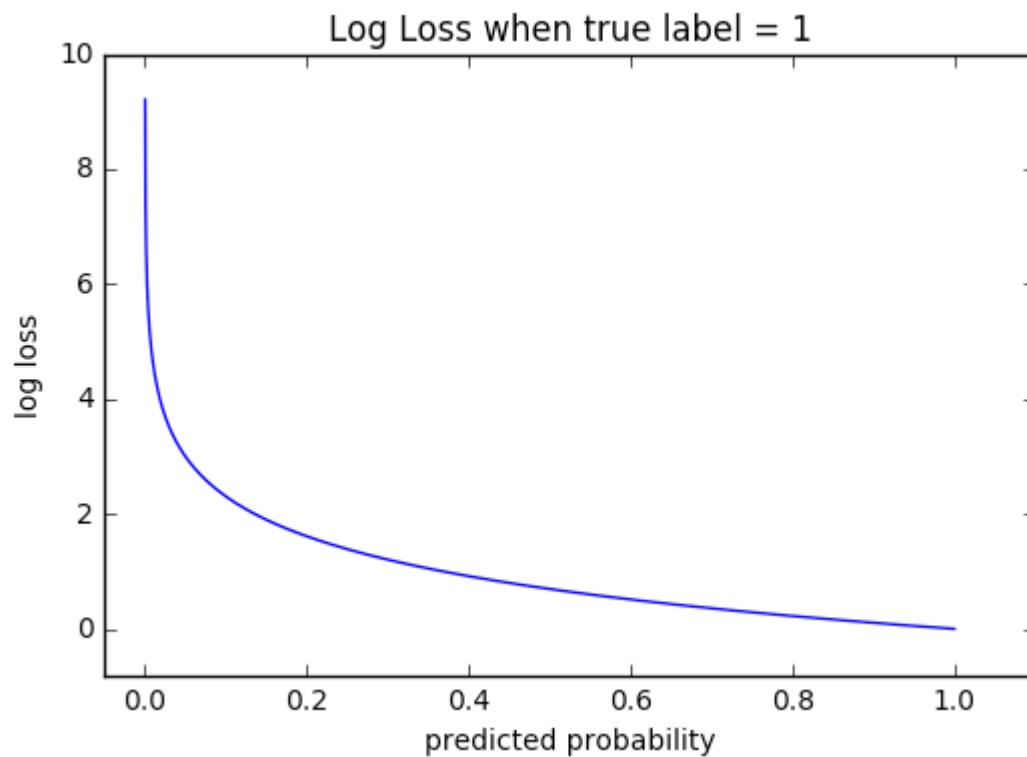
$$y = \frac{1}{1 + e^{-wx - b}}$$



$$\text{loss} = -(\check{y} \log(y) + (1 - \check{y}) \log(1 - y))$$

Cross entropy

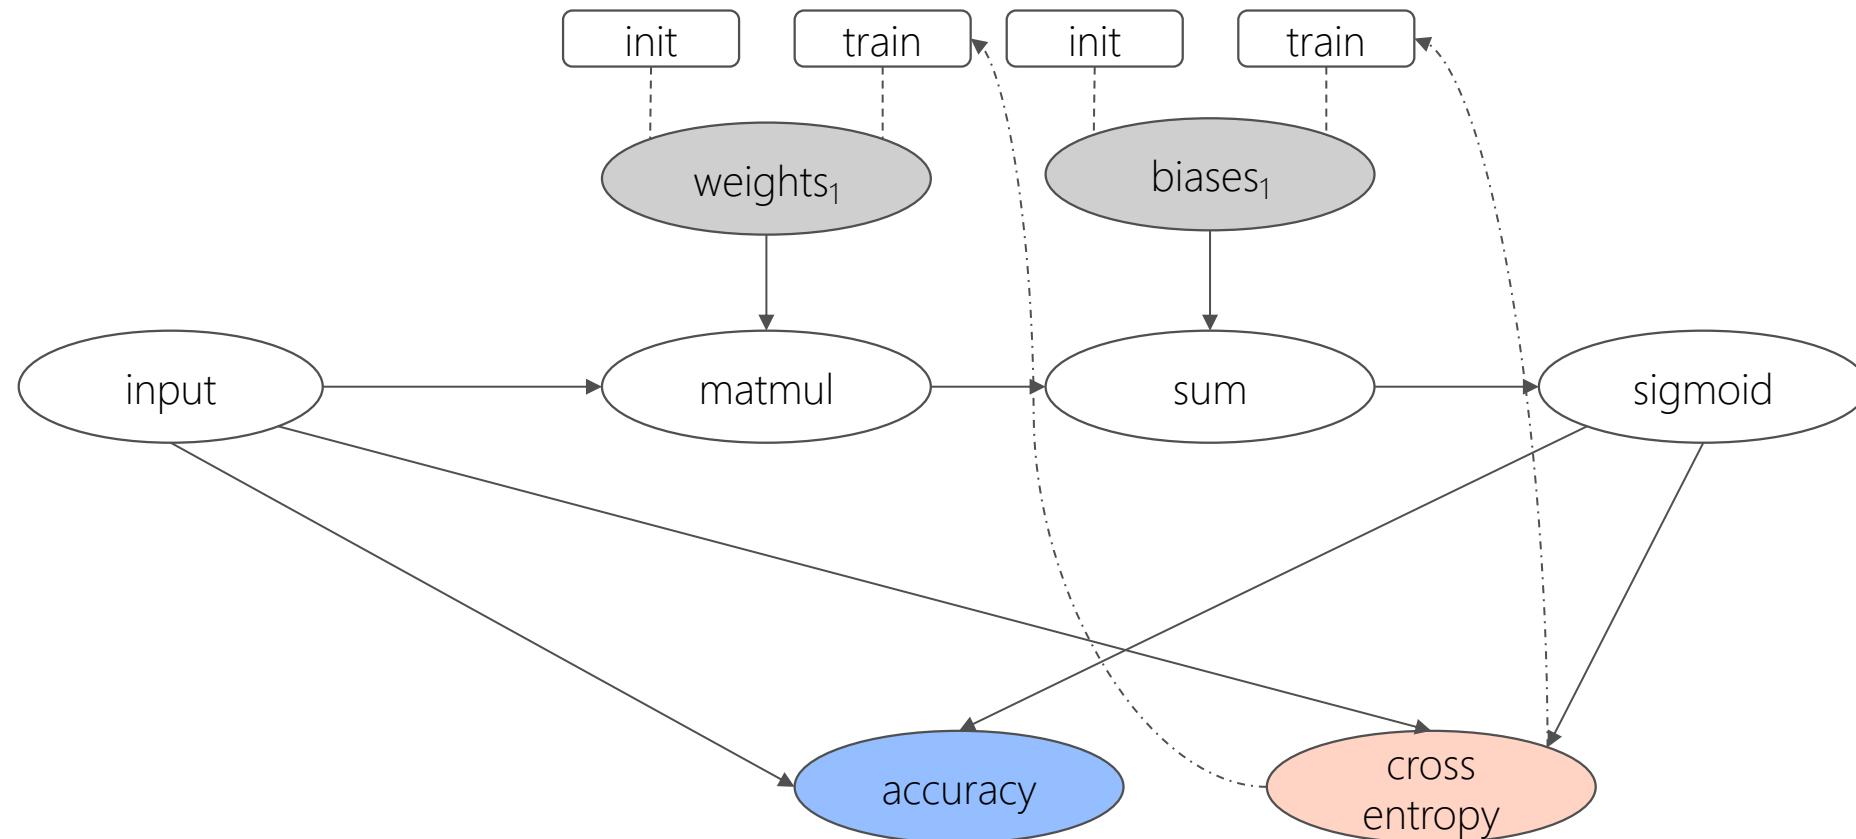
$$loss = -(\check{y} \log(y) + (1 - \check{y}) \log(1 - y))$$



Source: <http://bit.ly/2HhbJoG>

Step by step DAG: accuracy

$$y = \frac{1}{1 + e^{-wx - b}}$$

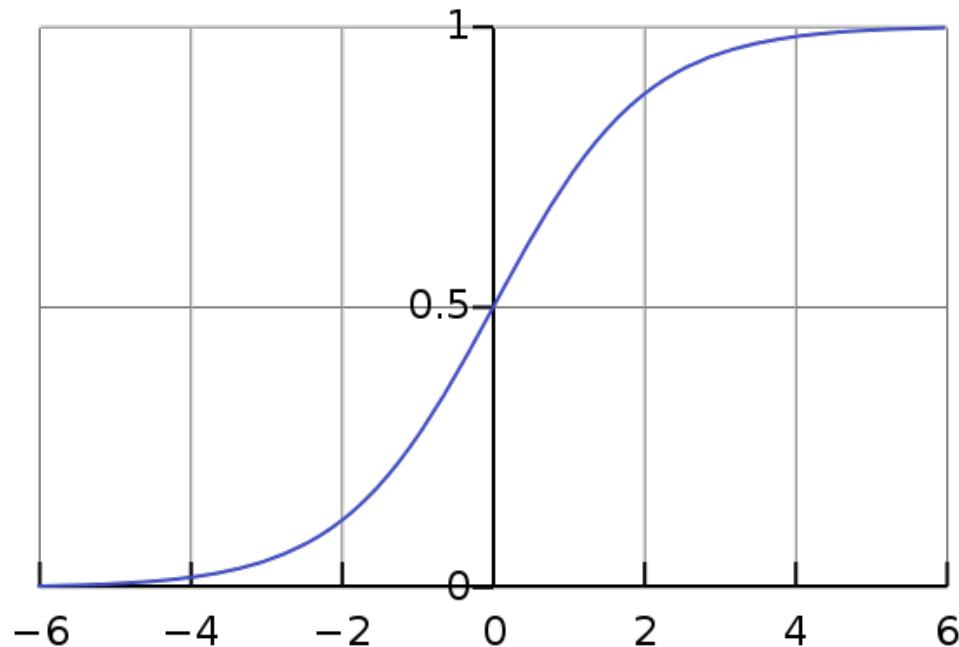


$$\text{loss} = -(\check{y} \log(y) + (1 - \check{y}) \log(1 - y))$$

Activation functions

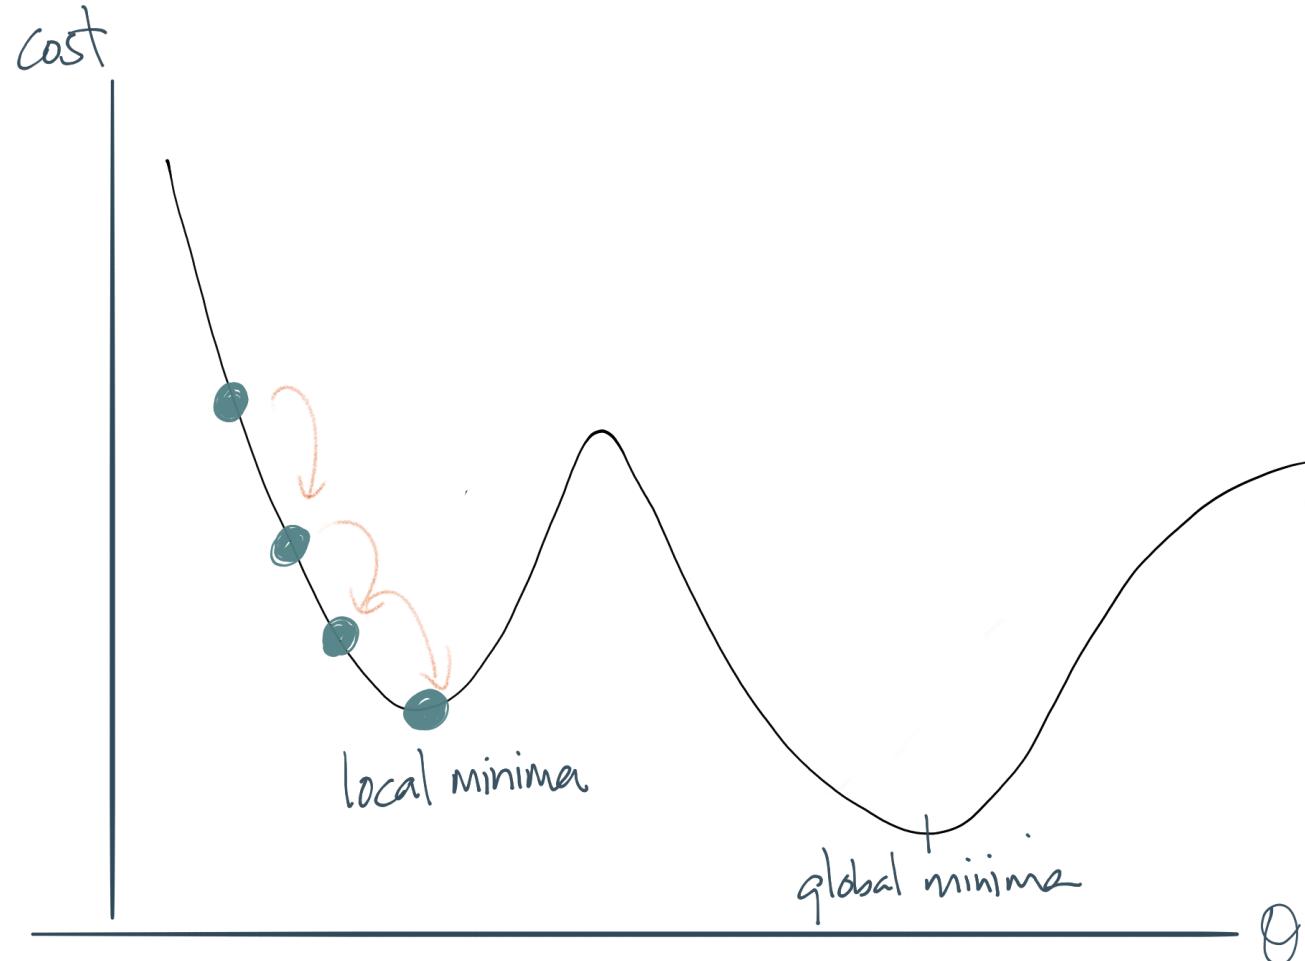
Sigmoid Function

$$f(x) = \sigma(z) \equiv \frac{1}{1 + e^{-z}} = \frac{1}{1 + \exp(-\sum_j w_j x_j - b)}$$



- Sigmoid non-linearity squashes real numbers between [0, 1]
- *Historically* nice interpretation of neuron firing rate (i.e. not firing at all (0) to fully-saturated firing (1)).

Gradient Descent

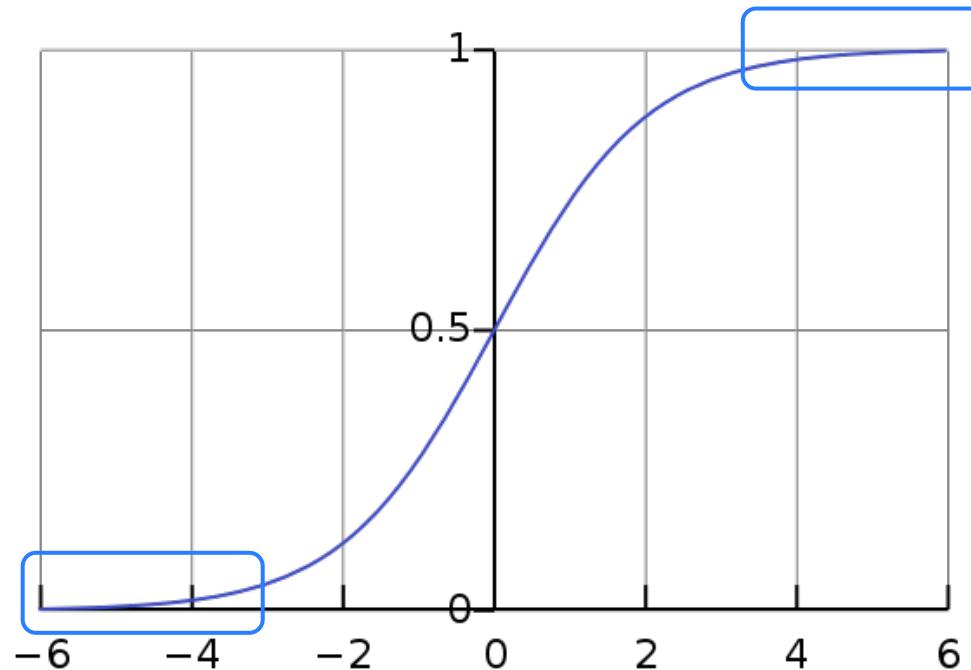


Cost function: what you want to minimize

Gradient descent: finding minimum of a function of multiple variables

Parameters like step size and learning rate are important to escape local minima

Sigmoid function (continued)

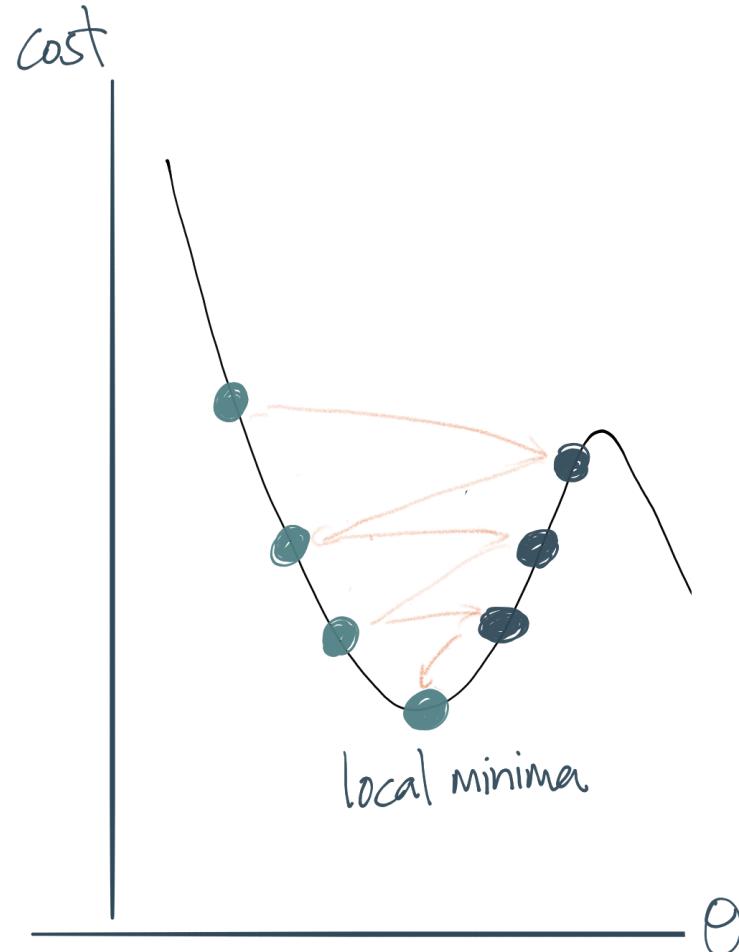


Source: <http://bit.ly/2GgMbGW>

Sigmoid function is not used as much b/c:

- Saturate and kill gradients: really large values are too close to 0 or 1, gradients are close to 0 (i.e. parallel to x-axis) thus stopping back propagation
- Output is not zero-centered: During gradient descent, if all values are positive then during back propagation the weights will become all positive or all negative creating zig zagging dynamics.

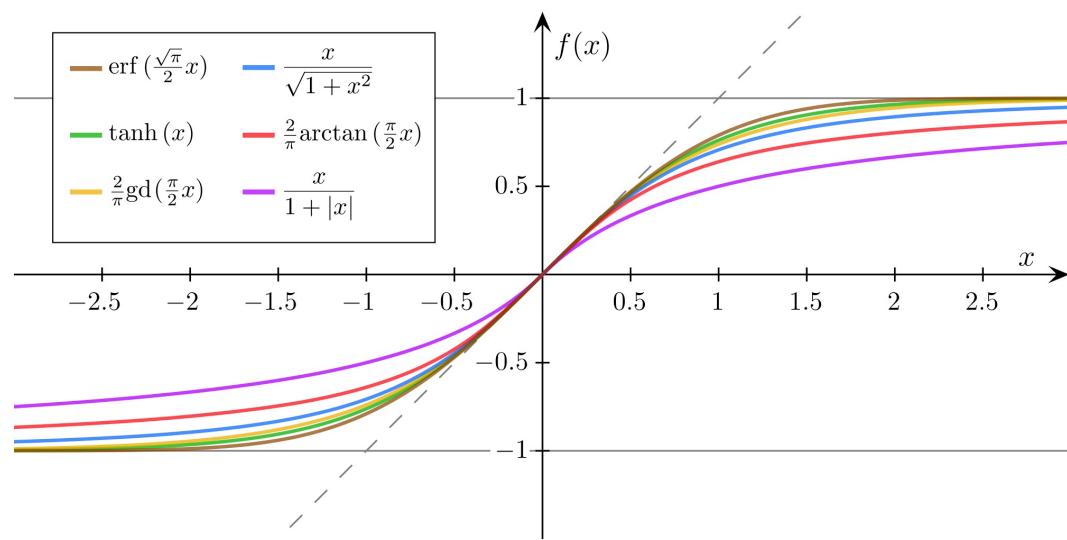
Sigmoid function (continued)



Sigmoid function is not used as much b/c:

- Saturate and kill gradients: really large values are too close to 0 or 1, gradients are close to 0 (i.e. parallel to x-axis) thus stopping back propagation
- Output is not zero-centered: During gradient descent, if all values are positive then during back propagation the weights will become all positive or all negative creating zig zagging dynamics.

Tanh function

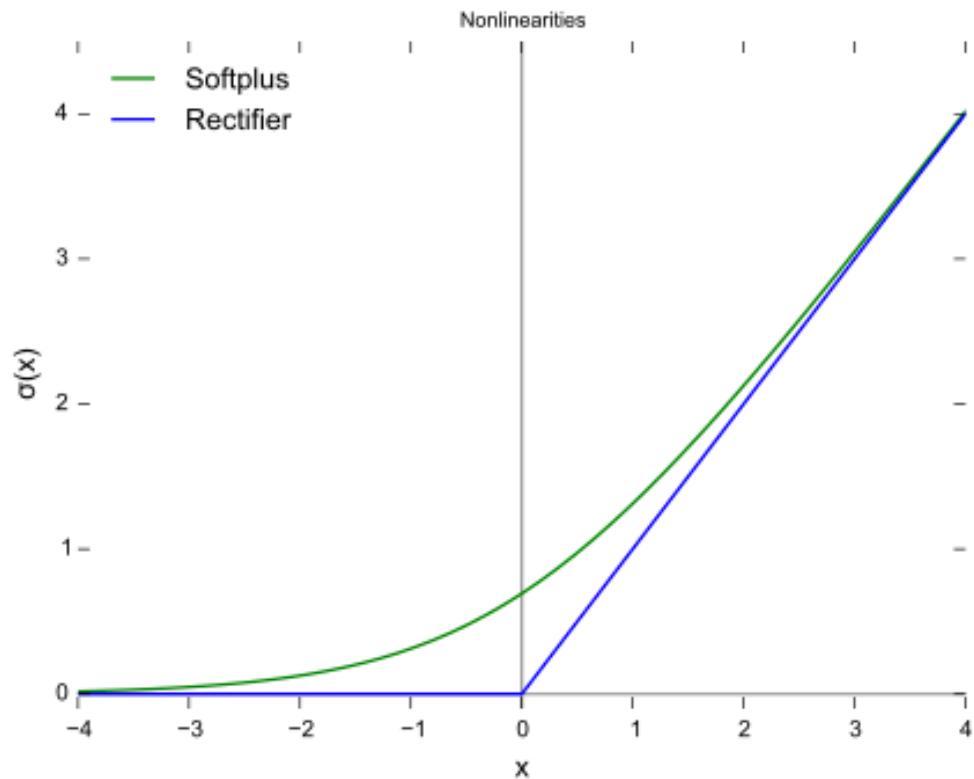


Source: <http://bit.ly/2C4y89z>

Tanh function squashes real numbers [-1, 1]

- Same problem as sigmoid that its activations saturate thus killing gradients.
- But it is zero-centered minimizing the zig zagging dynamics during gradient descent.
- Currently preferred *sigmoid nonlinearity*

ReLU: Rectifier Linear Unit



CC0, <http://bit.ly/2EwRnds>

ReLU's activation is thresholded at zero
 $f(z) = \max(0, u)$

- Quite popular over the last few years
- Speeds up SGD convergence
- It is easier to implement due to simpler mathematical functions
- Sensitive to high learning rate during training resulting in “dead” neurons (i.e. neurons that will not activate across the entire dataset).

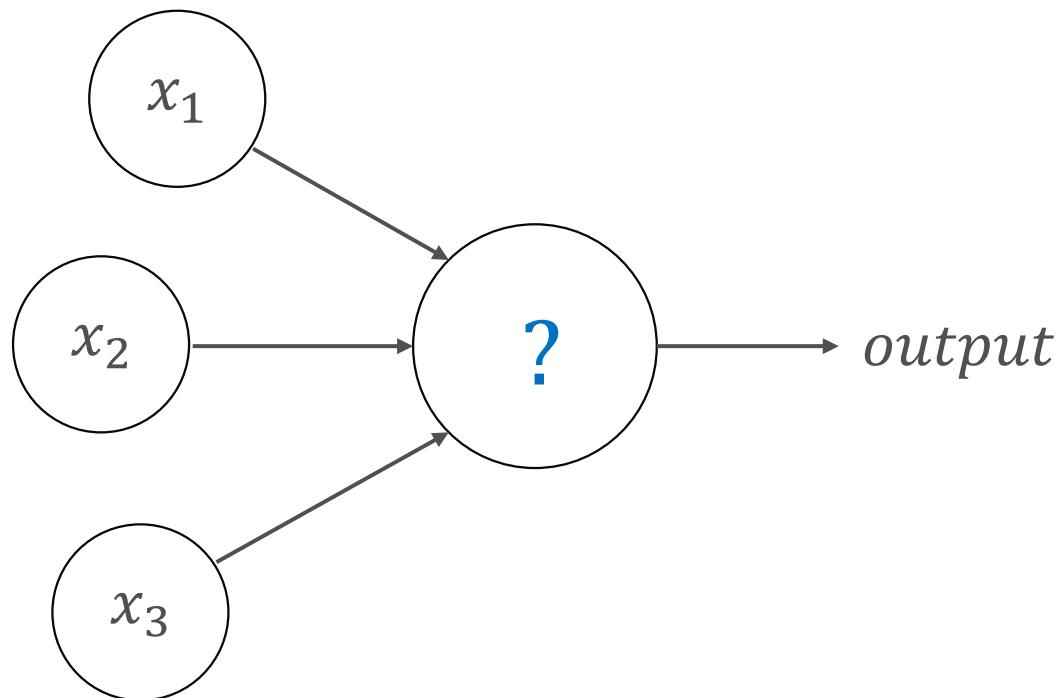
The Immense Power of Simple Structures

Resources:

- Andrej Karpathy's [CS231N Neural Networks](#)
- Steve Miller's [How to build a neural network](#)
- Michael Nielsen's [Neural Networks and Deep Learning, Chapter 1](#)

Perceptron

Simplified (binary) artificial neuron



*Do I snowboard at Whistler
this weekend?*

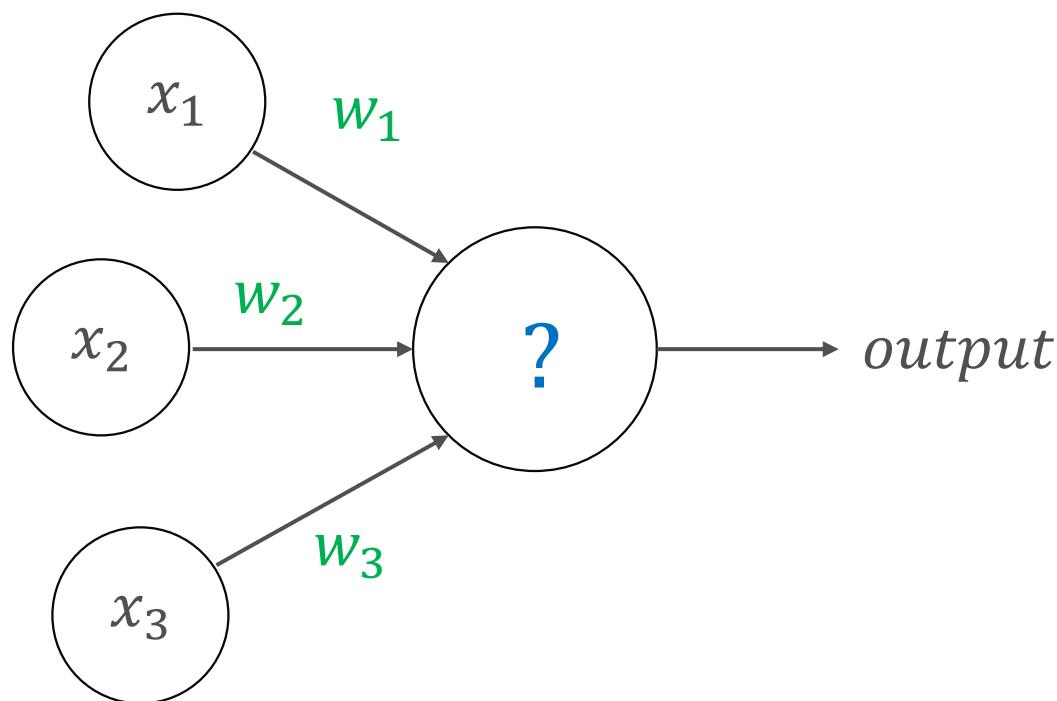
$x_1 \rightarrow$ Is the weather good?

$x_2 \rightarrow$ Is the powder good?

$x_3 \rightarrow$ Am I in the mood to drive?

Perceptron

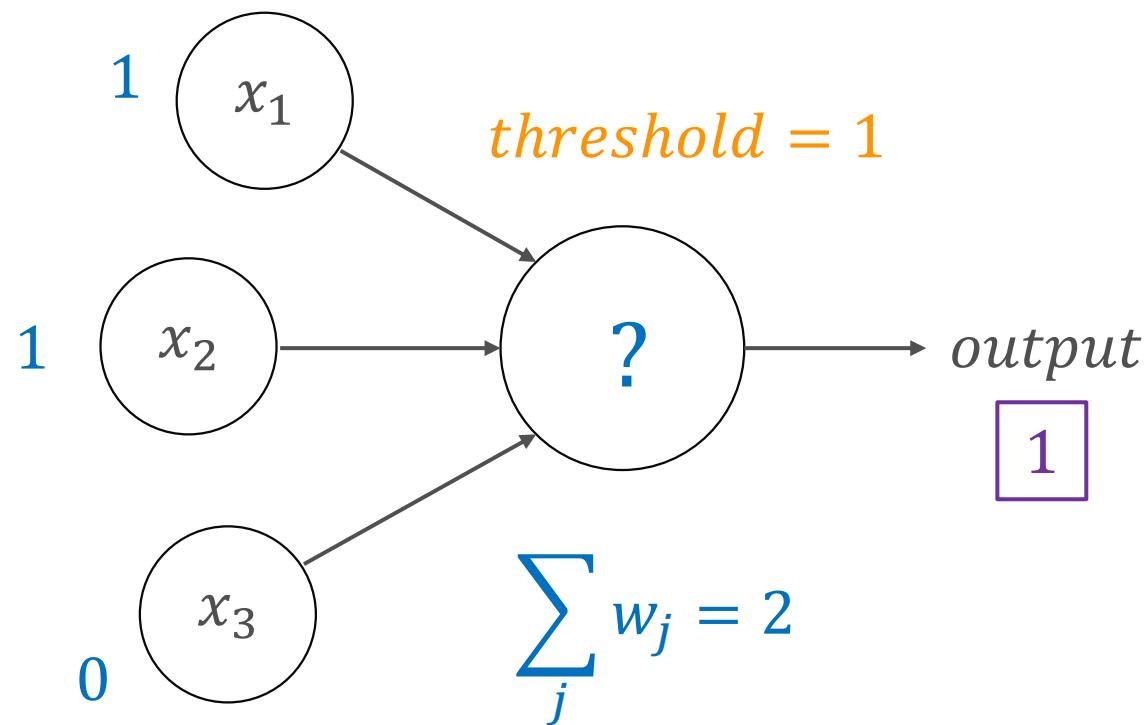
Simplified (binary) artificial neuron **with weights**



$$\text{output} = \begin{cases} 0, & \sum_{j=0}^n w_j x_j \leq \text{threshold} \\ 1, & \sum_{j=0}^n w_j x_j > \text{threshold} \end{cases}$$

Perceptron

Simplified (binary) artificial neuron; *no weights*



*Do I snowboard at Whistler
this weekend?*

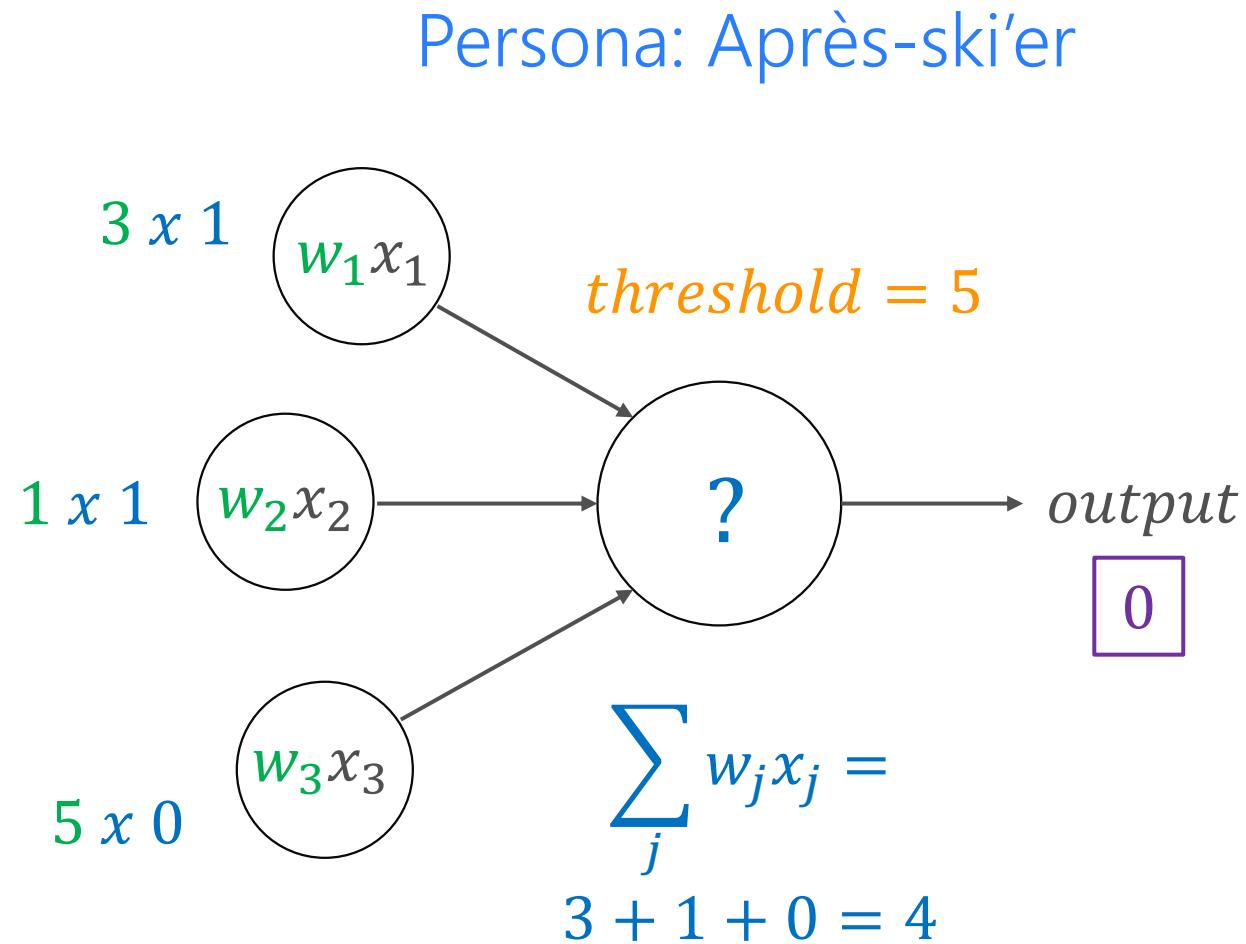
$x_1 = 1$ (*good weather*)

$x_2 = 1$ (*a lot of powder*)

$x_3 = 0$ (*driving sucks*)

Perceptron

Simplified (binary) artificial neuron; *add weights*



Do I snowboard at Whistler this weekend?

$x_1 = 1$ (*good weather*)

$w_1 = 3$

$x_2 = 1$ (*a lot of powder*)

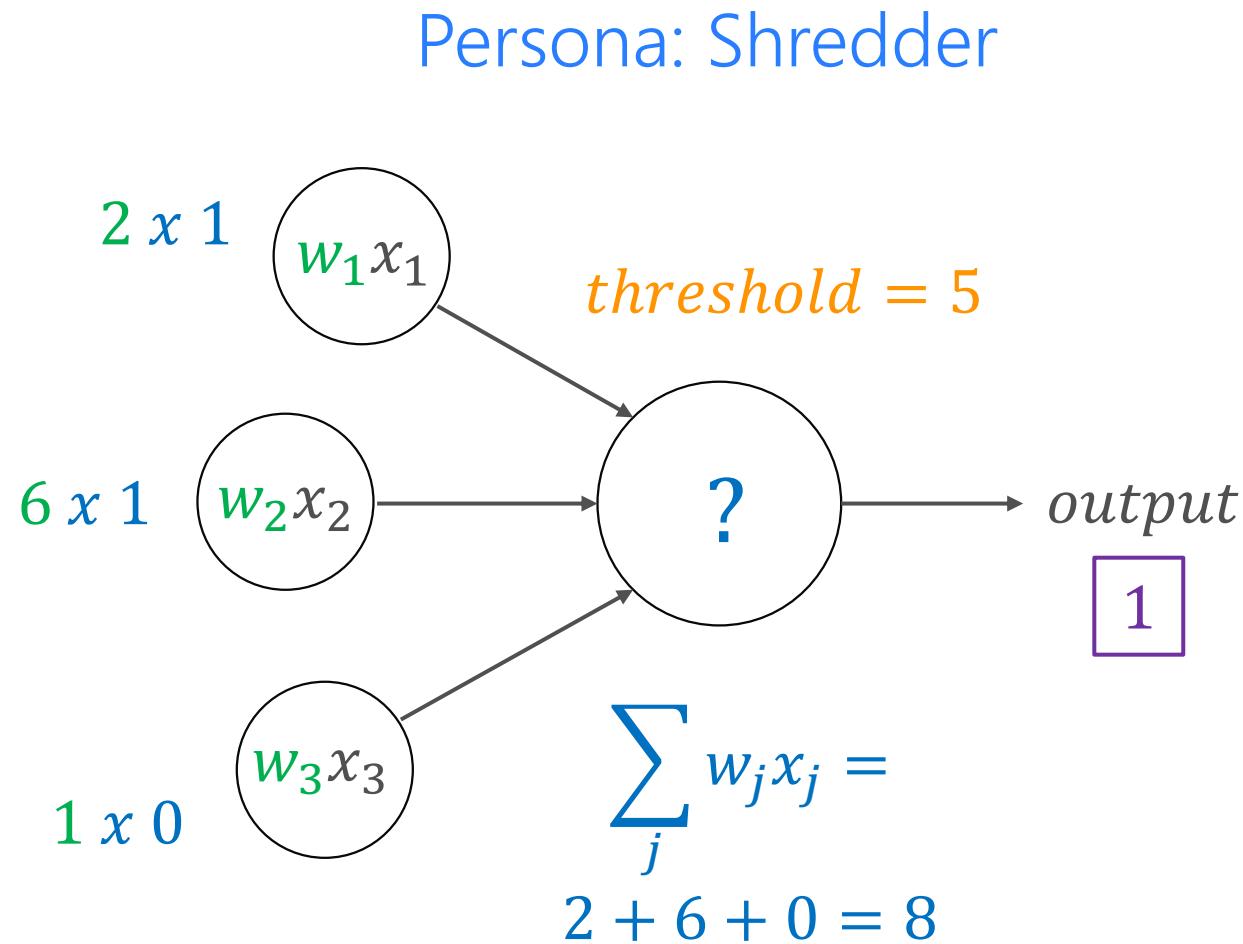
$w_2 = 1$

$x_3 = 0$ (*driving sucks*)

$w_3 = 5$

Perceptron

Simplified (binary) artificial neuron; *add weights*



Do I snowboard at Whistler this weekend?

$x_1 = 1$ (*good weather*)

$w_1 = 2$

$x_2 = 1$ (*a lot of powder*)

$w_2 = 6$

$x_3 = 0$ (*driving sucks*)

$w_3 = 1$

Introducing Bias

Perceptron needs to take into account of bias

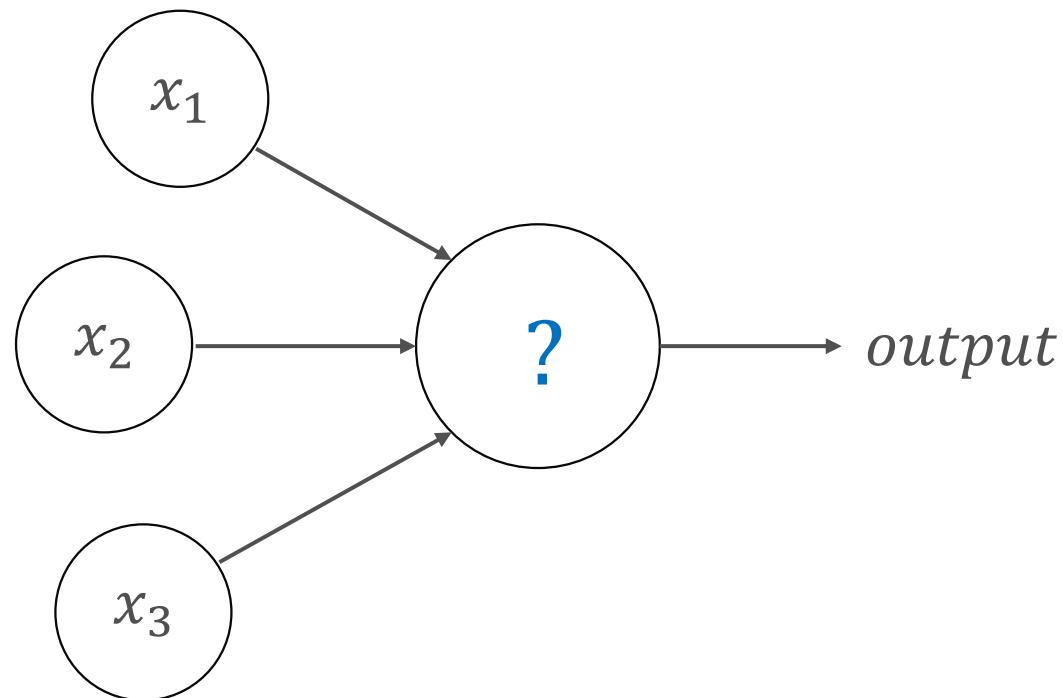
$$output = \begin{cases} 0, & wx + b \leq 0 \\ 1, & wx + b > 0 \end{cases}$$

where b is how easy it is to get the perceptron to fire

e.g. Shredder has a strong positive bias to go to Whistler
while Après-Ski'er bias is not as strong

Sigmoid Neuron

The more common artificial neuron

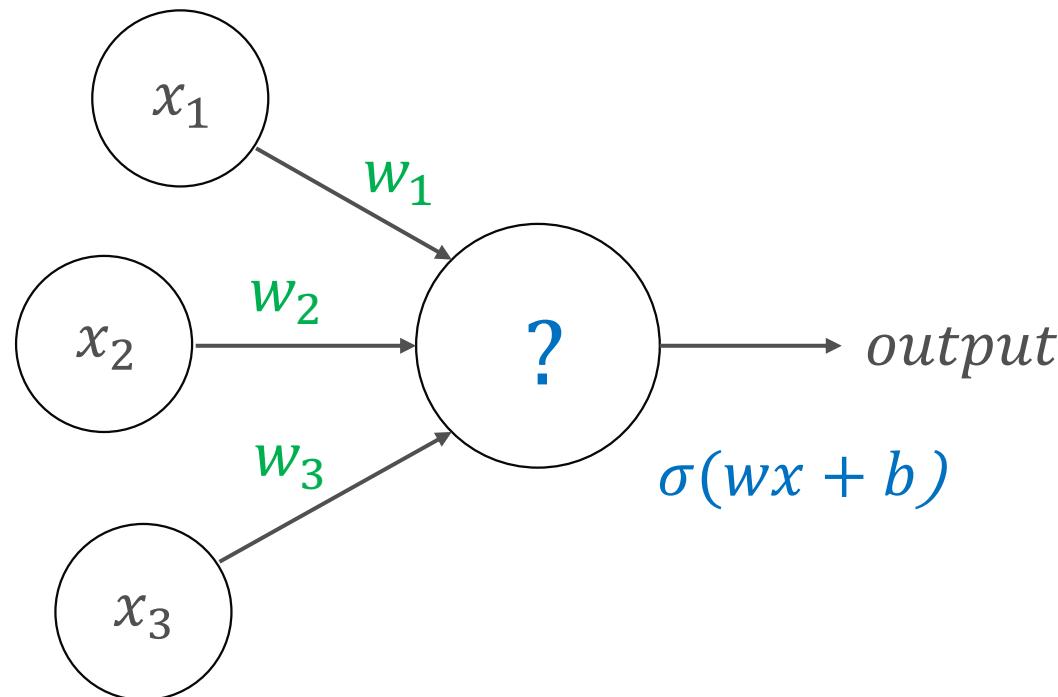


Instead of $[0, 1]$, now $(0\dots 1)$

Where output is defined by $\sigma(wx + b)$

Sigmoid Neuron

The more common artificial neuron

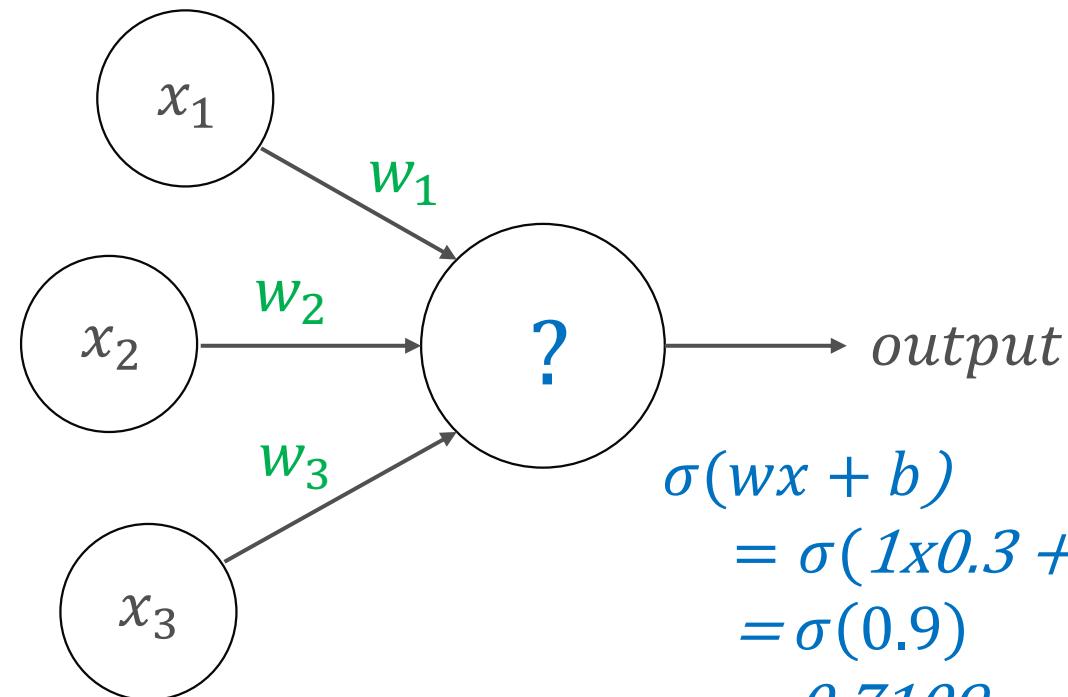


Instead of $[0, 1]$, now $(0\dots 1)$

Where output is defined by $\sigma(wx + b)$

Sigmoid Neuron

Persona: Shredder



$$\begin{aligned}\sigma(wx + b) \\ &= \sigma(1x0.3 + 1x0.6 + 0x 0.1) \\ &= \sigma(0.9) \\ &= 0.7109\end{aligned}$$

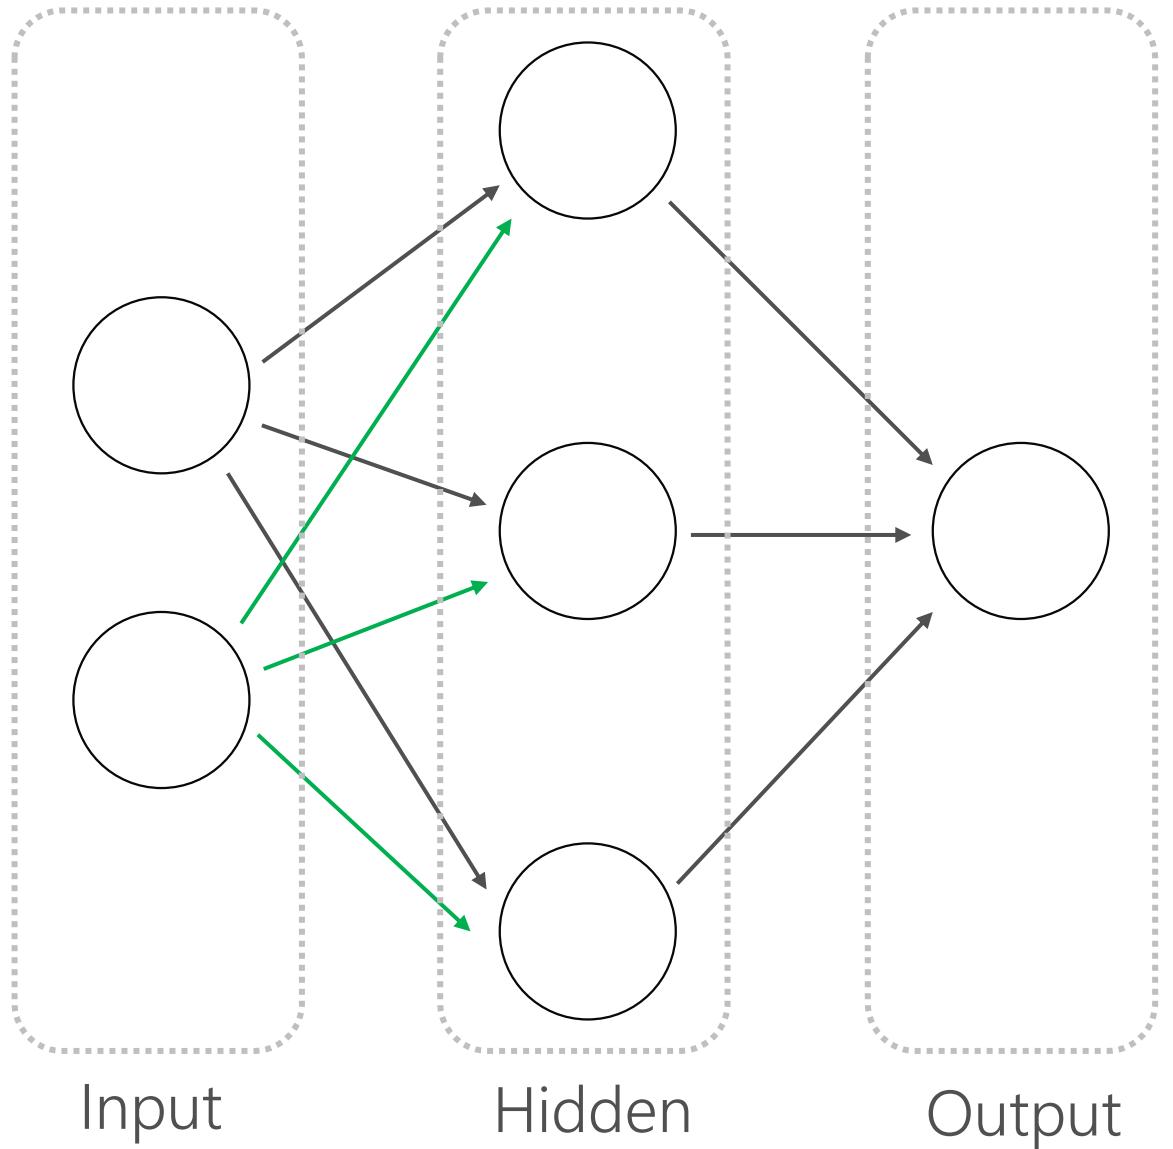
Do I snowboard at Whistler this weekend?

$x_1 = 1$ (*good weather*)
 $w_1 = 0.3$

$x_2 = 1$ (*a lot of powder*)
 $w_2 = 0.6$

$x_3 = 0$ (*driving sucks*)
 $w_3 = 0.1$

Simplified Two-Layer ANN



*Do I snowboard at Whistler
this weekend?*

$x_1 = \text{Après Ski'er}$

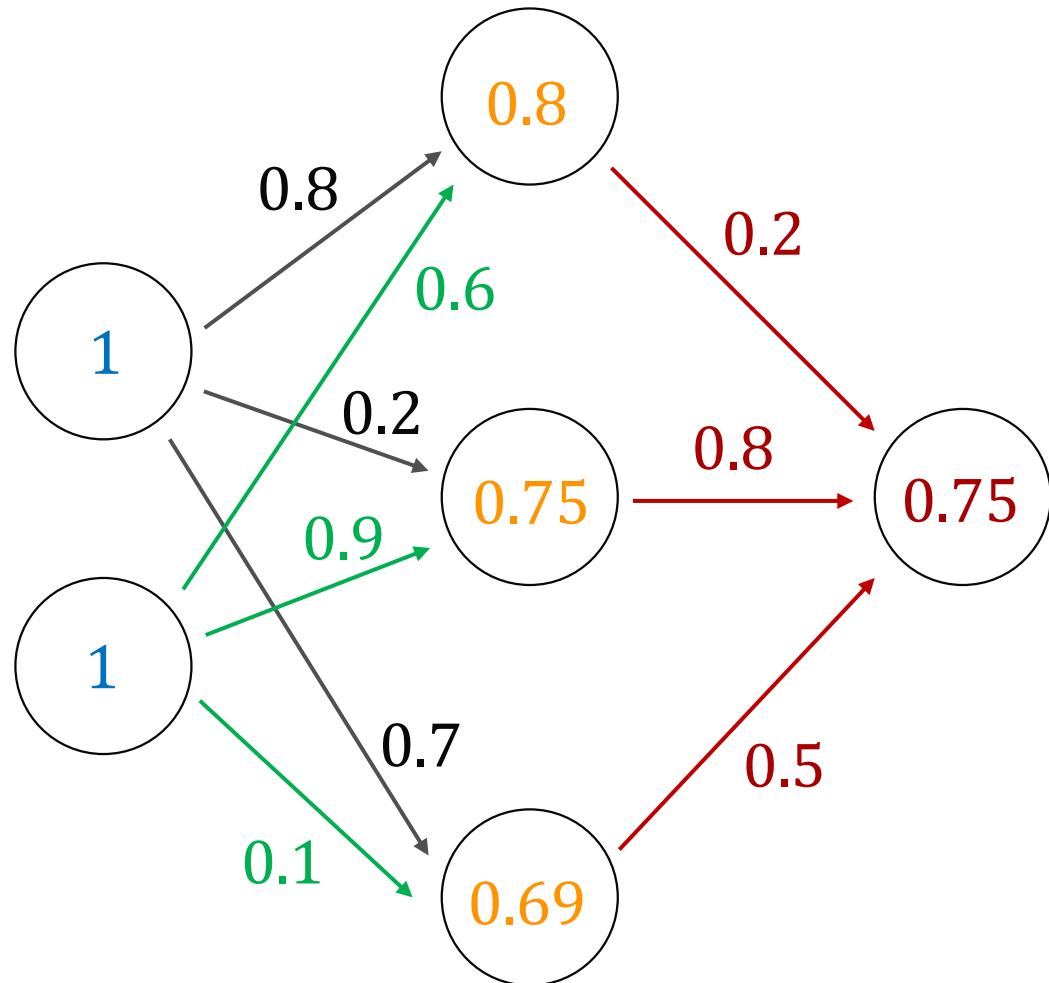
$x_2 = \text{Shredder}$

$h_1 = \text{weather}$

$h_2 = \text{powder}$

$h_3 = \text{driving}$

Simplified Two-Layer ANN



$$\begin{aligned} h_1 &= \sigma(1x0.8 + 1x0.6) = 0.80 \\ h_2 &= \sigma(1x0.2 + 1x0.9) = 0.75 \\ h_3 &= \sigma(1x0.7 + 1x0.1) = 0.69 \end{aligned}$$

output
= $\sigma(0.2x0.8 + 0.8x0.75 + 0.5x0.69)$
= $\sigma(1.105) = 0.75$

Great References

[Andrej Karpathy's ConvNetJS MNIST Demo](#)

[What is back propagation in neural networks?](#)

CS231n: Convolutional Neural Networks for Visual Recognition

[Syllabus and Slides](#) | [Course Notes](#) | [YouTube](#)

- With particular focus on [CS231n: Lecture 7: Convolution Neural Networks](#)

[Neural Networks and Deep Learning](#)

[TensorFlow](#)

Great References

[Deep Visualization Toolbox](#)

[Back Propagation with TensorFlow](#)

[TensorFrames: Google TensorFlow with Apache Spark](#)

[Integrating deep learning libraries with Apache Spark](#)

Upcoming webinar: [Build, Scale, and Deploy Deep Learning Pipelines with Ease](#)



Q&A